

Elements Of Data Science - F2025

Week 5: Intro to Machine Learning Models

9/30/2025

TODOs

- **Readings:**
 - Recommended: Take a look through https://scikit-learn.org/stable/supervised_learning.html
 - Reference: PML Chapter Chap 3
- **Quiz 5:** Due Oct 14st, 11:59pm ET
 - *Be sure to indicate on which page each question output is displayed*
- **HW1:** Due Oct 14th, 11:59pm ET
 - *Be sure to indicate on which page each question output is displayed*
- **HW2:** out the week after midterm

Today

- Intro to Machine Learning Models
- Various types of ML
- Linear models

Next class

- One Vs. Rest For Multiclass/Multilabel Classification
- Distance Based: kNN
- Tree Based: Decision Tree
- Ensembles: Bagging, Boosting, Stacking

Questions?

Environment Setup

Environment Setup

In [1]:

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5
6 sns.set_style('darkgrid')
7
8 %matplotlib inline
```

Modeling and ML

- What is a (ML) Model?
 - Specification of a mathematical (or probabilistic) relationship between different variables.
- What is Machine Learning?
 - Creating and using models that learn from data.

Questions for Models

```
In [2]: 1 ! head ../data/wine_dataset.csv
```

```
alcohol,malic_acid,ash,alcalinity_of_ash,magnesium,total_phenols,flavanoids,nonflavanoid_phenols,proanthocyanins,color_intensit  
y,hue,od280/od315_of_diluted_wines,proline,class  
14.23,1.71,2.43,15.6,127.0,2.8,3.06,0.28,2.29,5.64,1.04,3.92,1065.0,0  
13.2,1.78,2.14,11.2,100.0,2.65,2.76,0.26,1.28,4.38,1.05,3.4,1050.0,0  
13.16,2.36,2.67,18.6,101.0,2.8,3.24,0.3,2.81,5.68,1.03,3.17,1185.0,0  
14.37,1.95,2.5,16.8,113.0,3.85,3.49,0.24,2.18,7.8,0.86,3.45,1480.0,0  
13.24,2.59,2.87,21.0,118.0,2.8,2.69,0.39,1.82,4.32,1.04,2.93,735.0,0  
14.2,1.76,2.45,15.2,112.0,3.27,3.39,0.34,1.97,6.75,1.05,2.85,1450.0,0  
14.39,1.87,2.45,14.6,96.0,2.5,2.52,0.3,1.98,5.25,1.02,3.58,1290.0,0  
14.06,2.15,2.61,17.6,121.0,2.6,2.51,0.31,1.25,5.05,1.06,3.58,1295.0,0  
14.83,1.64,2.17,14.0,97.0,2.8,2.98,0.29,1.98,5.2,1.08,2.85,1045.0,0
```

Questions for Models

```
In [2]: 1 ! head ../data/wine_dataset.csv
```

```
alcohol,malic_acid,ash,alcalinity_of_ash,magnesium,total_phenols,flavanoids,nonflavanoid_phenols,proanthocyanins,color_intensit  
y,hue,od280/od315_of_diluted_wines,proline,class  
14.23,1.71,2.43,15.6,127.0,2.8,3.06,0.28,2.29,5.64,1.04,3.92,1065.0,0  
13.2,1.78,2.14,11.2,100.0,2.65,2.76,0.26,1.28,4.38,1.05,3.4,1050.0,0  
13.16,2.36,2.67,18.6,101.0,2.8,3.24,0.3,2.81,5.68,1.03,3.17,1185.0,0  
14.37,1.95,2.5,16.8,113.0,3.85,3.49,0.24,2.18,7.8,0.86,3.45,1480.0,0  
13.24,2.59,2.87,21.0,118.0,2.8,2.69,0.39,1.82,4.32,1.04,2.93,735.0,0  
14.2,1.76,2.45,15.2,112.0,3.27,3.39,0.34,1.97,6.75,1.05,2.85,1450.0,0  
14.39,1.87,2.45,14.6,96.0,2.5,2.52,0.3,1.98,5.25,1.02,3.58,1290.0,0  
14.06,2.15,2.61,17.6,121.0,2.6,2.51,0.31,1.25,5.05,1.06,3.58,1295.0,0  
14.83,1.64,2.17,14.0,97.0,2.8,2.98,0.29,1.98,5.2,1.08,2.85,1045.0,0
```

```
In [3]: 1 df_wine = pd.read_csv('../data/wine_dataset.csv',usecols=['alcohol','ash','proline','hue','class'])  
2 df_wine.sample(4,random_state=1)
```

Out[3]:

	alcohol	ash	hue	proline	class
161	13.69	2.54	0.96	680.0	2
117	12.42	2.19	1.06	345.0	1
19	13.64	2.56	0.96	845.0	0
69	12.21	1.75	1.28	718.0	1

Data Vocabulary for ML

Data Vocabulary for ML

- **Classification:** Can we predict *categorical target/label* "class" from the other columns?
- **Regression:** Can we predict *numeric target* "hue" from the other columns?
- **Feature Selection:** What are the *important features* when predicting "hue"?
- **Interpretation:** Can a model tell us about how the features and target/label interact?
- **Clustering:** Do the observations group together in feature space?

Data Vocabulary for ML

- **Classification:** Can we predict *categorical target/label* "class" from the other columns?
- **Regression:** Can we predict *numeric target* "hue" from the other columns?
- **Feature Selection:** What are the *important features* when predicting "hue"?
- **Interpretation:** Can a model tell us about how the features and target/label interact?
- **Clustering:** Do the observations group together in feature space?
- X , features, attributes, independent/exogenous/explanatory variables
 - Ex: alcohol, trip_distance, company_industry
- y , target, label, outcome, dependent/endogenous/response variables
 - Ex: class, hue, tip_amount, stock_price
- $f(X) \rightarrow y$, Model that maps features X to target y

Data Vocabulary for ML

- **Classification:** Can we predict *categorical target/label* "class" from the other columns?
- **Regression:** Can we predict *numeric target* "hue" from the other columns?
- **Feature Selection:** What are the *important features* when predicting "hue"?
- **Interpretation:** Can a model tell us about how the features and target/label interact?
- **Clustering:** Do the observations group together in feature space?
- X , features, attributes, independent/exogenous/explanatory variables
 - Ex: alcohol, trip_distance, company_industry
- y , target, label, outcome, dependent/endogenous/response variables
 - Ex: class, hue, tip_amount, stock_price
- $f(X) \rightarrow y$, Model that maps features X to target y

```
In [4]: 1 df_wine.sample(5,random_state=1)
```

Out[4]:

	alcohol	ash	hue	proline	class
161	13.69	2.54	0.96	680.0	2
117	12.42	2.19	1.06	345.0	1

Variations of ML Tasks

Variations of ML Tasks

- **Supervised vs Unsupervised**
 - is there a target/label?

Variations of ML Tasks

- **Supervised vs Unsupervised**
 - is there a target/label?
- **Regression vs Classification**
 - is the target numeric or categorical?

Variations of ML Tasks

- **Supervised vs Unsupervised**
 - is there a target/label?
- **Regression vs Classification**
 - is the target numeric or categorical?
- **Prediction vs Interpretation**
 - generate predictions or understand interactions?

Variations of ML Tasks

- **Supervised vs Unsupervised**
 - is there a target/label?
- **Regression vs Classification**
 - is the target numeric or categorical?
- **Prediction vs Interpretation**
 - generate predictions or understand interactions?
- **Model Families**
 - Linear, Tree, Distance, Probability, Neural Net, Ensemble, ...

Supervised vs Unsupervised vs Reinforcement Learning

- Is there a target, y ?
 - Yes
 - **Supervised Learning:** Data consists of (X, y) pairs
 - Uses: Classification, Regression
 - Ex: What is the relationship between length of ride and tip amount?
 - No
 - **Unsupervised Learning:** Data consists only of (X)
 - Uses: Clustering, Topic Modeling, etc.
 - Ex: Are there any clusters in length of ride?
 - Eventually?
 - **Reinforcement Learning**
 - After a series of predictions (path) get a reward from a reward function
 - Ex: Poker player

Other Learning Paradigms

Other Learning Paradigms

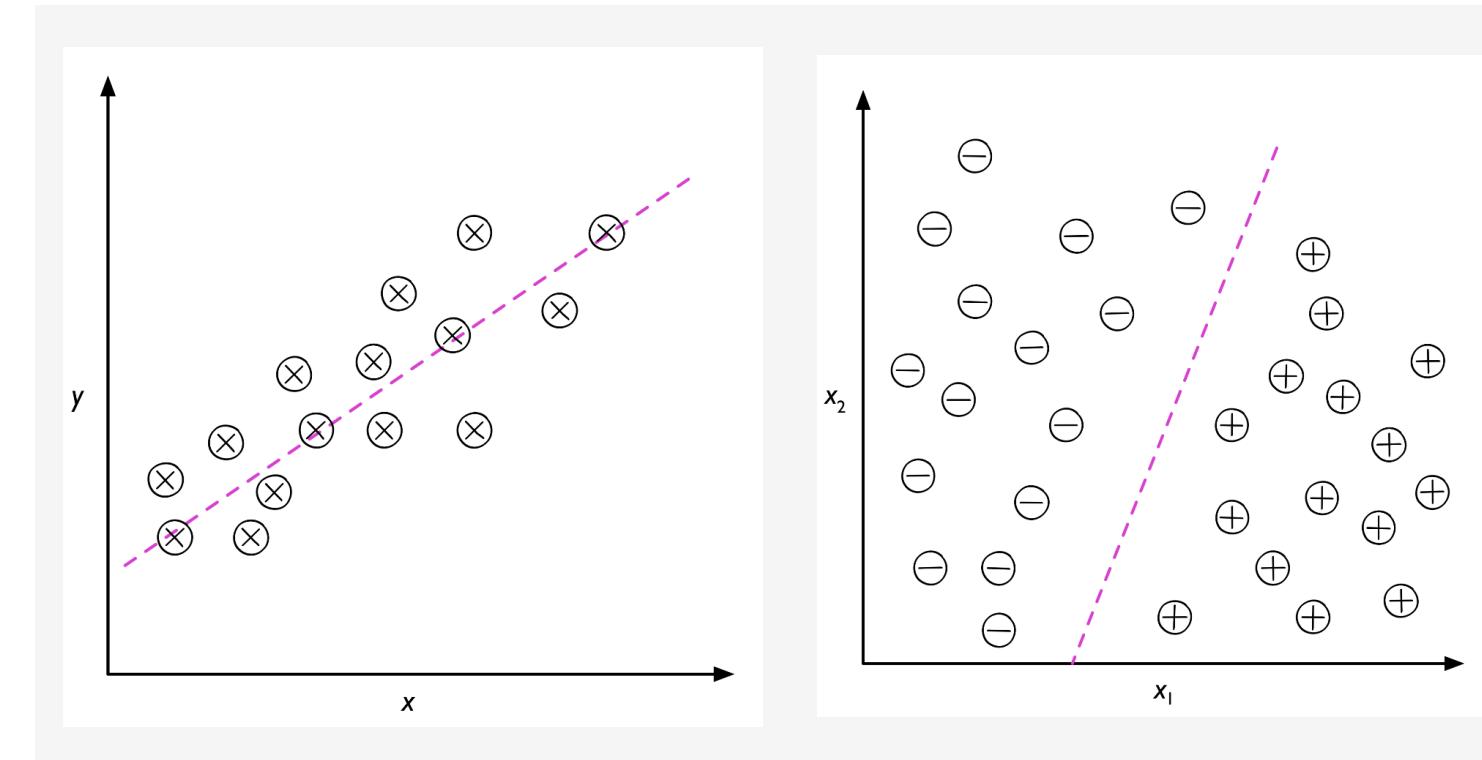
- Do we have a mix of labeled and unlabeled?
 - **Semi-Supervised Learning**
 - Can we use structure of unlabeled data along with labeled?

Other Learning Paradigms

- Do we have a mix of labeled and unlabeled?
 - **Semi-Supervised Learning**
 - Can we use structure of unlabeled data along with labeled?
- Will we continue getting new data?
 - **Online Learning**
 - Is there an oracle (ground truth) we can consult?
 - Can we select which points to make predictions on?

Supervised Learning: Regression vs Classification

- **Regression** -> predict a numeric value
 - Ex: tip_amount, stock_price, wine_hue
- **Classification** -> predict a discrete class/category
 - Ex: class of wine, face/not face, object labels in image
- Note: convert a regression problem into classification with binning/thresholding



Prediction vs Interpretation

- Do we care more about **prediction**: how good are our predictions?
 - Ex: For a given taxi trip, what will the tip size likely be?
 - Ex: For a given loan, will there be a default?
- Or, do we care more about **interpretation**: understanding how X relates to y ?
 - Ex: What happens to tip size as taxi trip length increases?
 - Ex: What is the relationship between debt and loan default?

Model Families for Supervised Learning

- Linear
 - Simple/Multiple Linear Regression
 - Logistic Regression (for Classification)
 - Support Vector Machines
 - Perceptron
- Tree Based
 - Decision Tree
- Distance Based
 - K-Nearest Neighbor

Model Families for Supervised Learning Continued

- Probability
 - Naive Bayes
 - Bayes Net
- Ensemble
 - Random Forest
 - Gradient Boosted Trees
 - Stacking
- Network
 - Multi-layer Perceptron
 - Deep Neural-Networks
 - Convolutional Neural Nets
 - Recurrent Neural Nets
 - Transformers

Example: Regression with a Linear Model

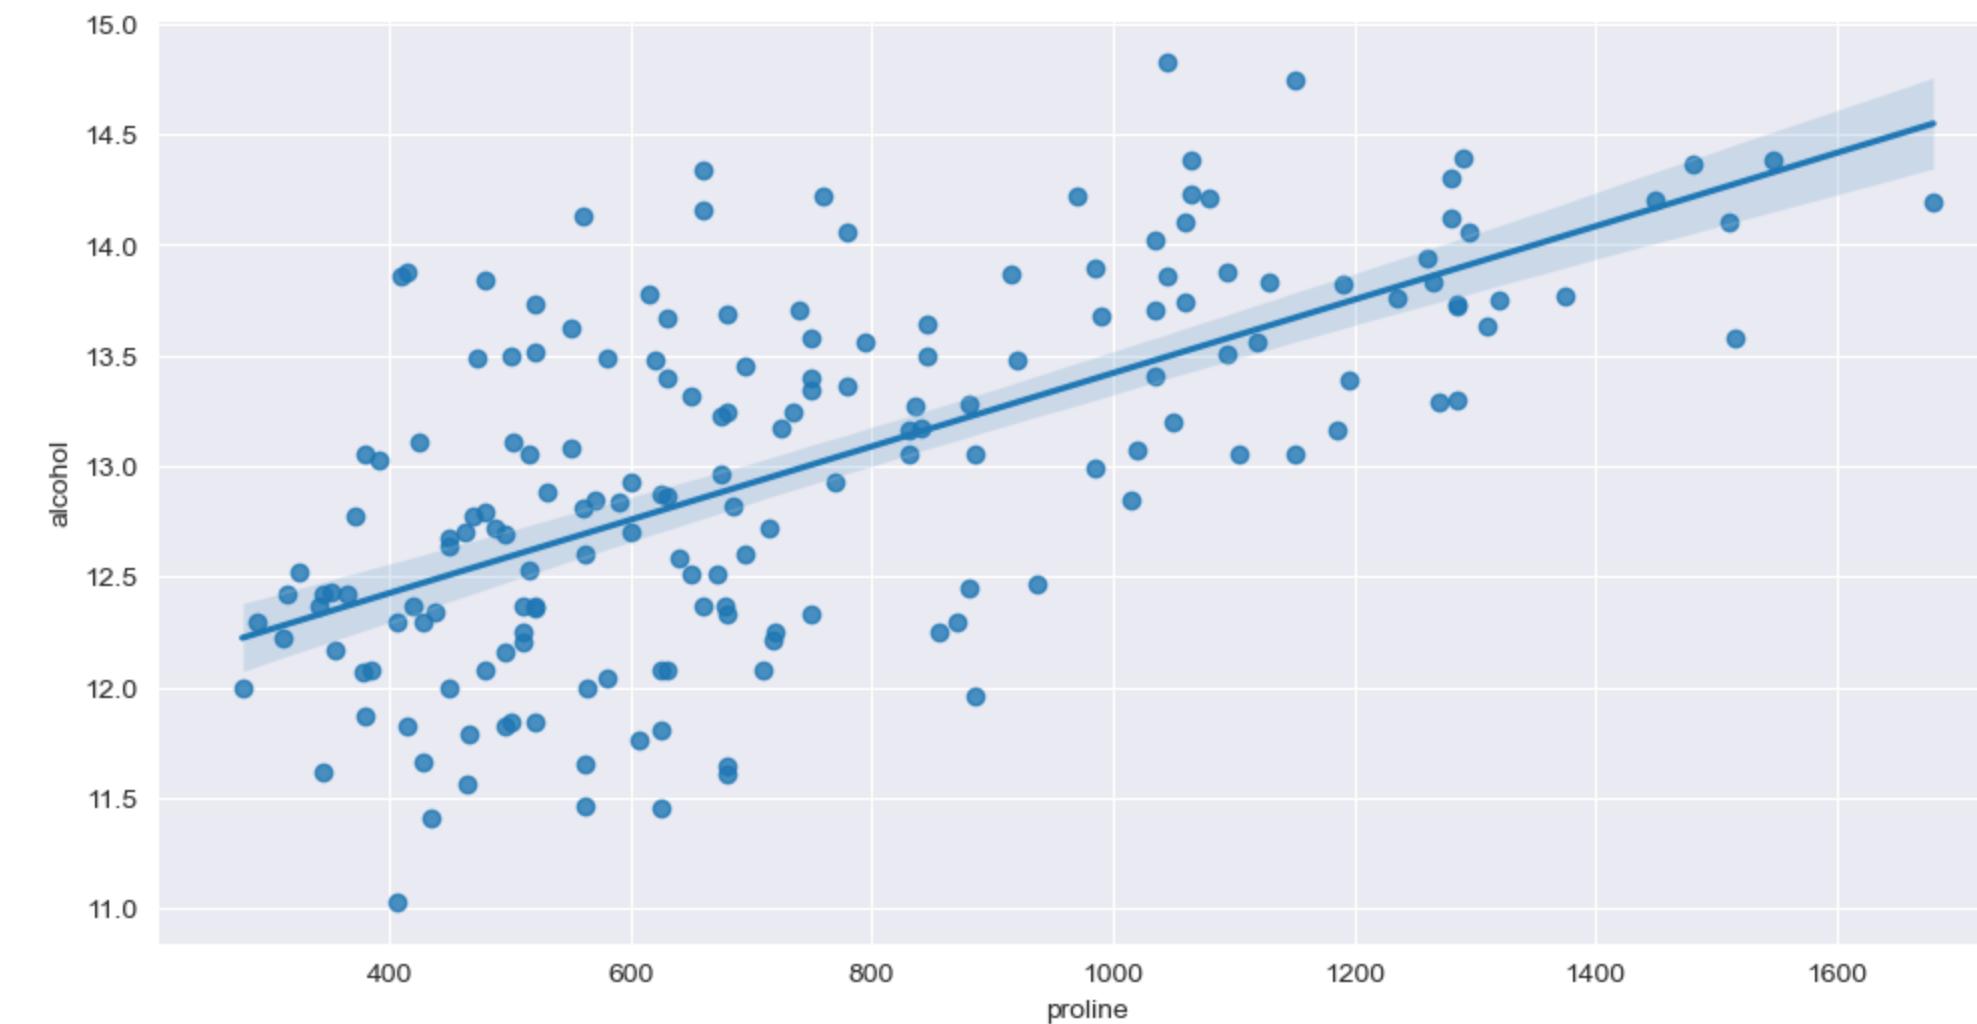
What is the relationship between 'proline' (an amino-acid) and 'alcohol' in wine?

Example: Regression with a Linear Model

What is the relationship between 'proline' (an amino-acid) and 'alcohol' in wine?

In [5]:

```
1 fig,ax = plt.subplots(1,1,figsize=(12,6))
2 sns.regplot(x='proline', y='alcohol', data=df_wine);
```



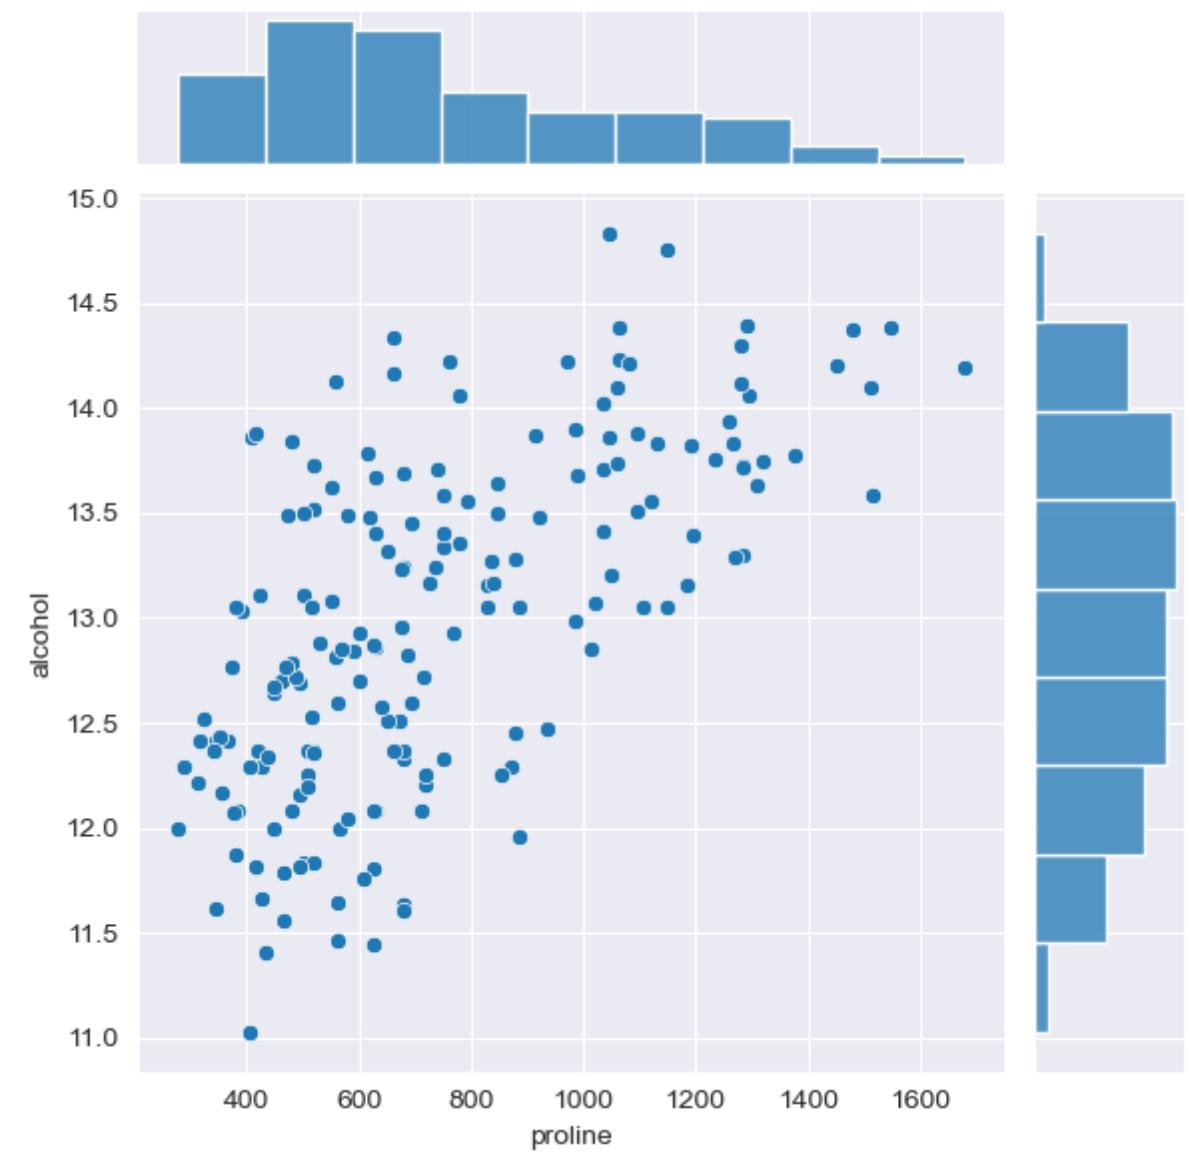
Aside: Correlation

Question: are proline and alcohol correlated?

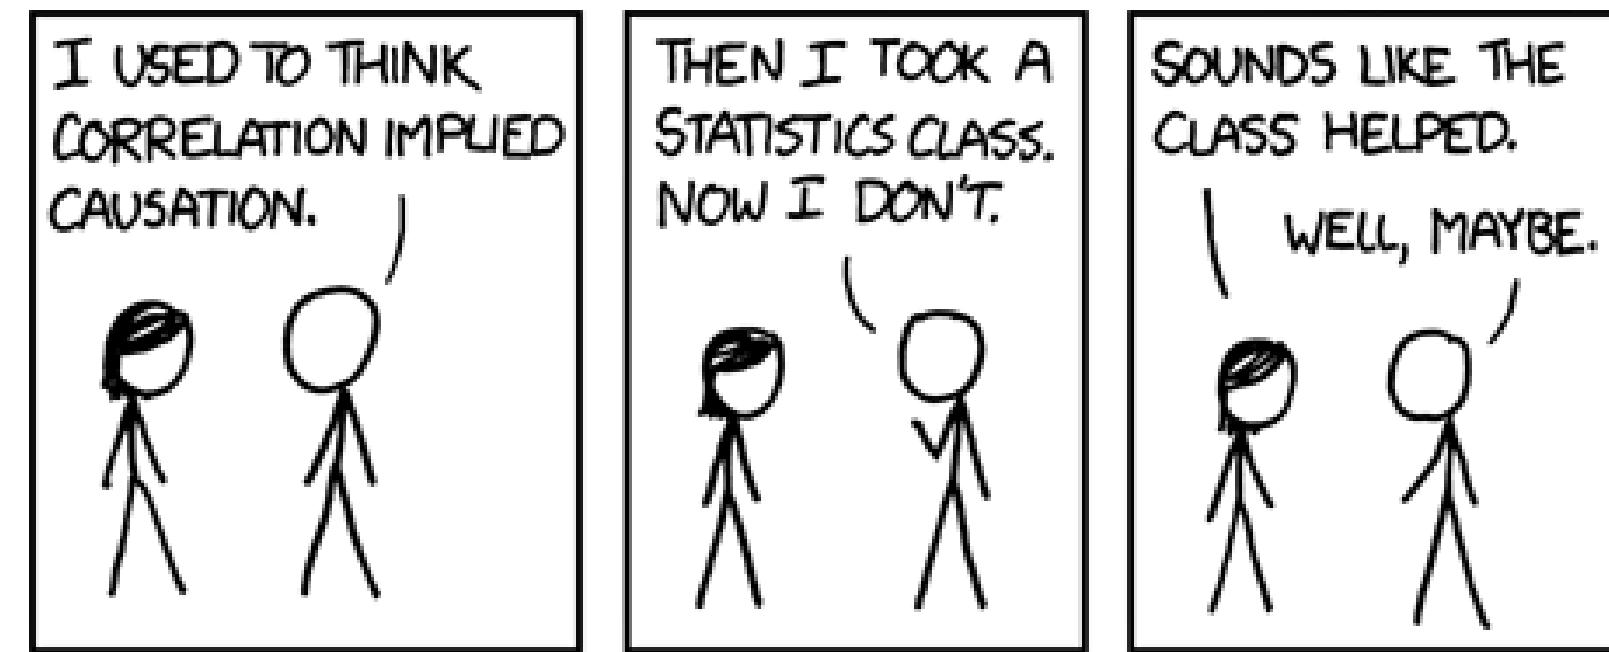
Aside: Correlation

Question: are proline and alcohol correlated?

```
In [6]: 1 sns.jointplot(x='proline',y='alcohol',data=df_wine);
```



Obligitory Correlation vs. Causation

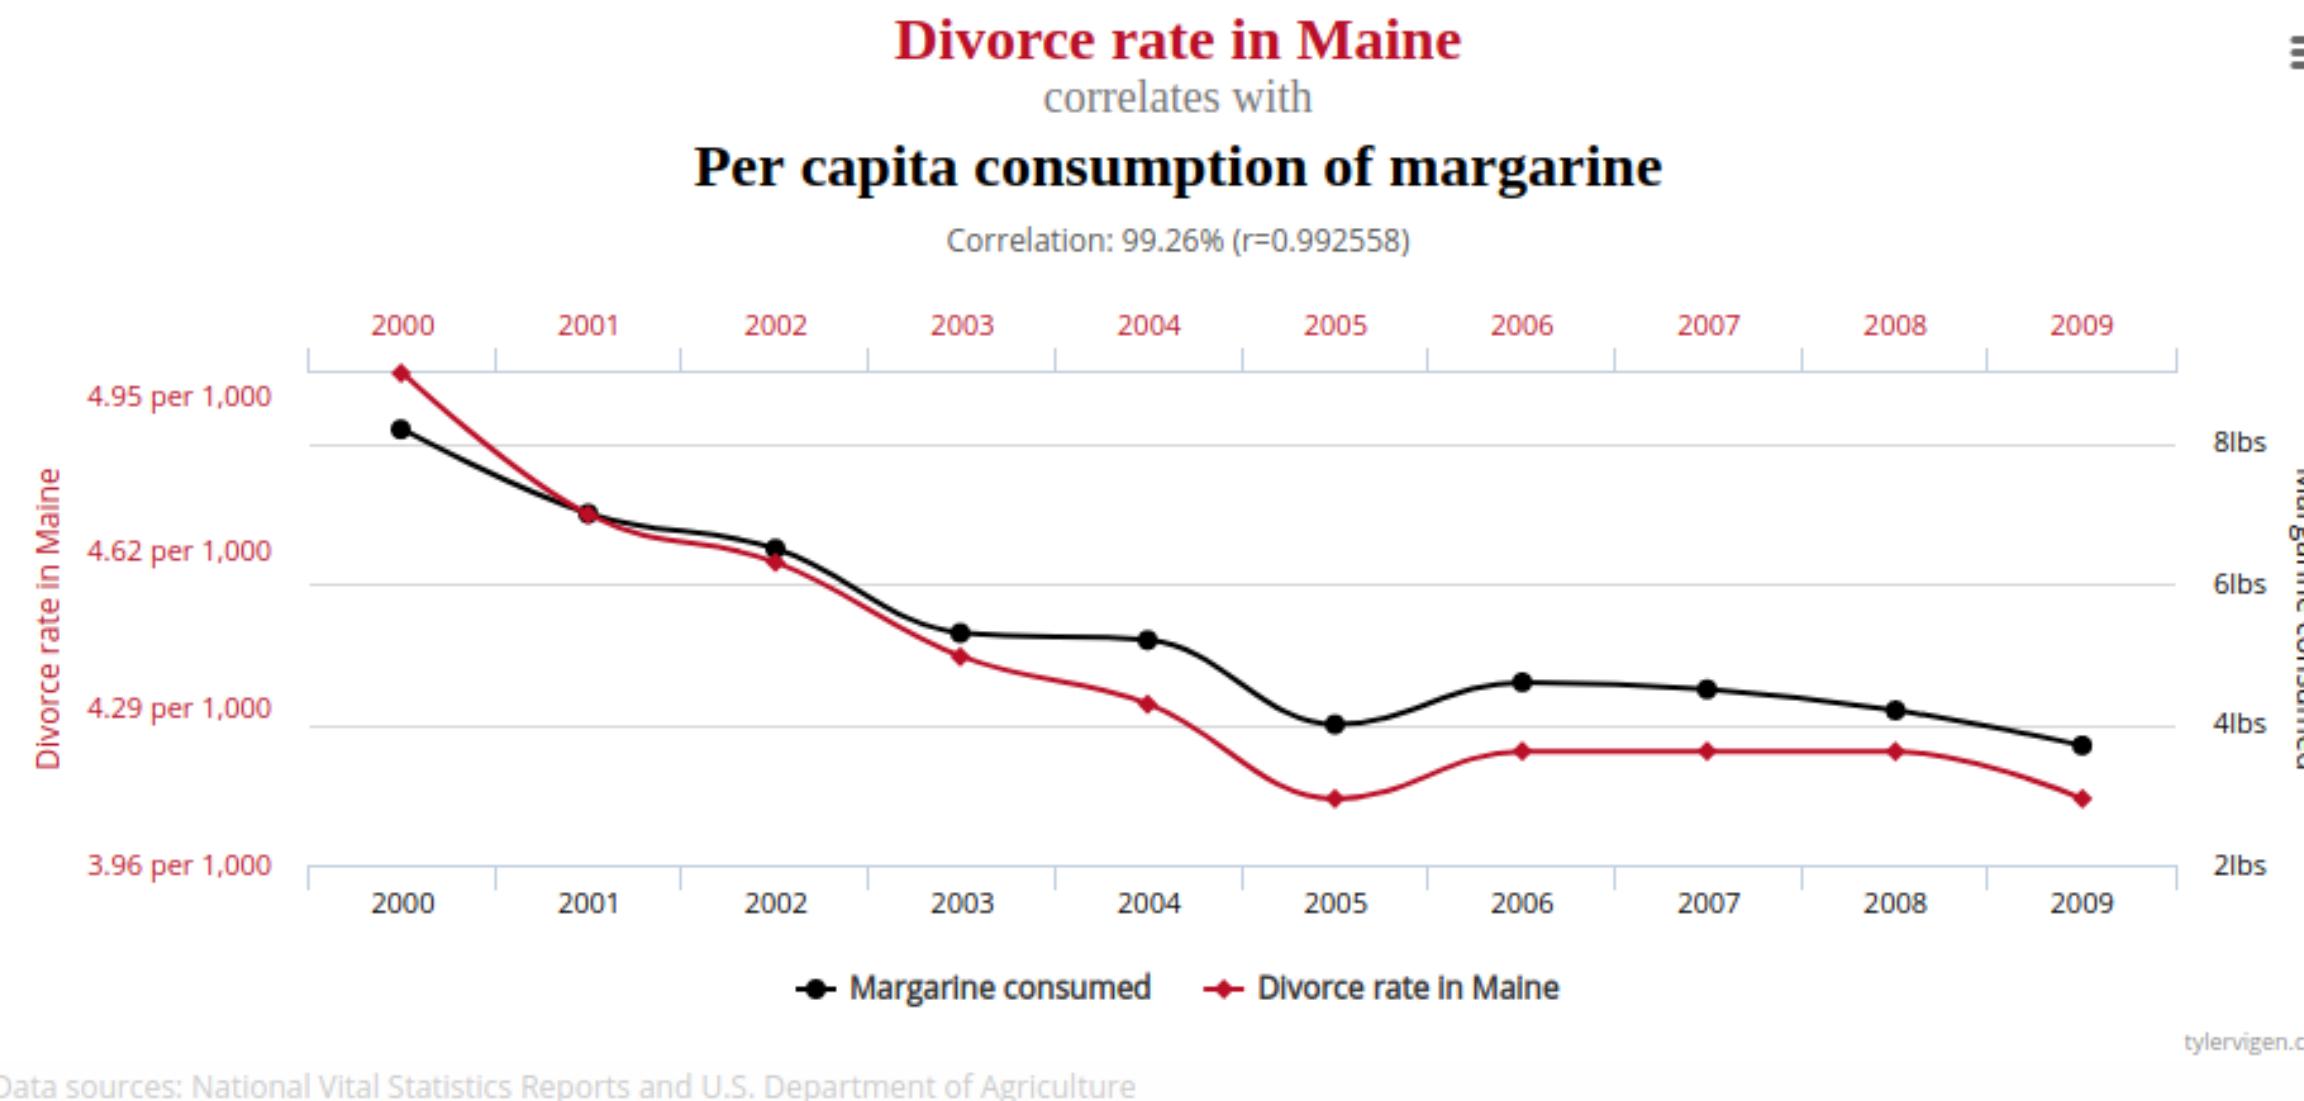


<https://imgs.xkcd.com/comics/correlation.png>

- Correlation does not imply causation!
- Look to Causal Inference
 - controlled experiment
 - control for confounding variables

Spurious Correlation

- Also, look hard enough and you'll find correlation.
 - See spurious correlations for examples



Aside: Correlation

- Could calculate Pearson Correlation Coefficient
- Assumes normally distributed data! (which is not true here)
 - On the Effects of Non-Normality on the Distribution of the Sample Product-Moment Correlation Coefficient

```
In [7]: 1 print(f"pearsonr : {df_wine[['proline','alcohol']].corr().iloc[0,1].round(2)}")  
pearsonr : 0.64
```

Aside: Correlation

- Could calculate Pearson Correlation Coefficient
- Assumes normally distributed data! (which is not true here)
 - On the Effects of Non-Normality on the Distribution of the Sample Product-Moment Correlation Coefficient

```
In [7]: 1 print(f"pearsonr : {df_wine[['proline','alcohol']].corr().iloc[0,1].round(2)}")  
pearsonr : 0.64
```

```
In [8]: 1 from scipy.stats import pearsonr  
2 r,p = pearsonr(df_wine.proline,df_wine.alcohol)  
3 print(f'r: {r:.2f}, p: {p:.2f}')  
  
r: 0.64, p: 0.00
```

Aside: Correlation

- Could calculate Pearson Correlation Coefficient
- Assumes normally distributed data! (which is not true here)
 - On the Effects of Non-Normality on the Distribution of the Sample Product-Moment Correlation Coefficient

```
In [7]: 1 print(f"pearsonr : {df_wine[['proline','alcohol']].corr().iloc[0,1].round(2)}")  
pearsonr : 0.64
```

```
In [8]: 1 from scipy.stats import pearsonr  
2 r,p = pearsonr(df_wine.proline,df_wine.alcohol)  
3 print(f'r: {r:.2f}, p: {p:.2f}')  
  
r: 0.64, p: 0.00
```

- We know that as proline goes up alcohol goes up, but by how much?

Python Modeling Libraries

Prediction - scikit-learn



Interpretation - scikit-learn and statsmodels



Additional Tools - mlxtend



Aside: MLxtend and conda-forge

- **MLxtend:** (machine learning extensions) is a Python library of useful tools for the day-to-day data science tasks.



- Recall: **Conda-Forge:** A community-led collection of recipes, build infrastructure and distributions for the conda package manager.



Current Timeline (subject to change)

- Midterm (In two weeks, in class)
- HW2 (Due Nov 11th)
- HW3 (Due Nov 25th)
- HW4 (Due Dec 9nd)
- Final (Dec 16th)

Midterm

- In-class
- 20 Multiple Choice questions
- 5 short answers

Topics to concentrate on

- conda virtual environment
- assert
- numpy indexing
- list comprehension
- pandas
- series and dataframes
- indexing
- Boolean indexing
- Quantiles
- Correlation
- Visualization

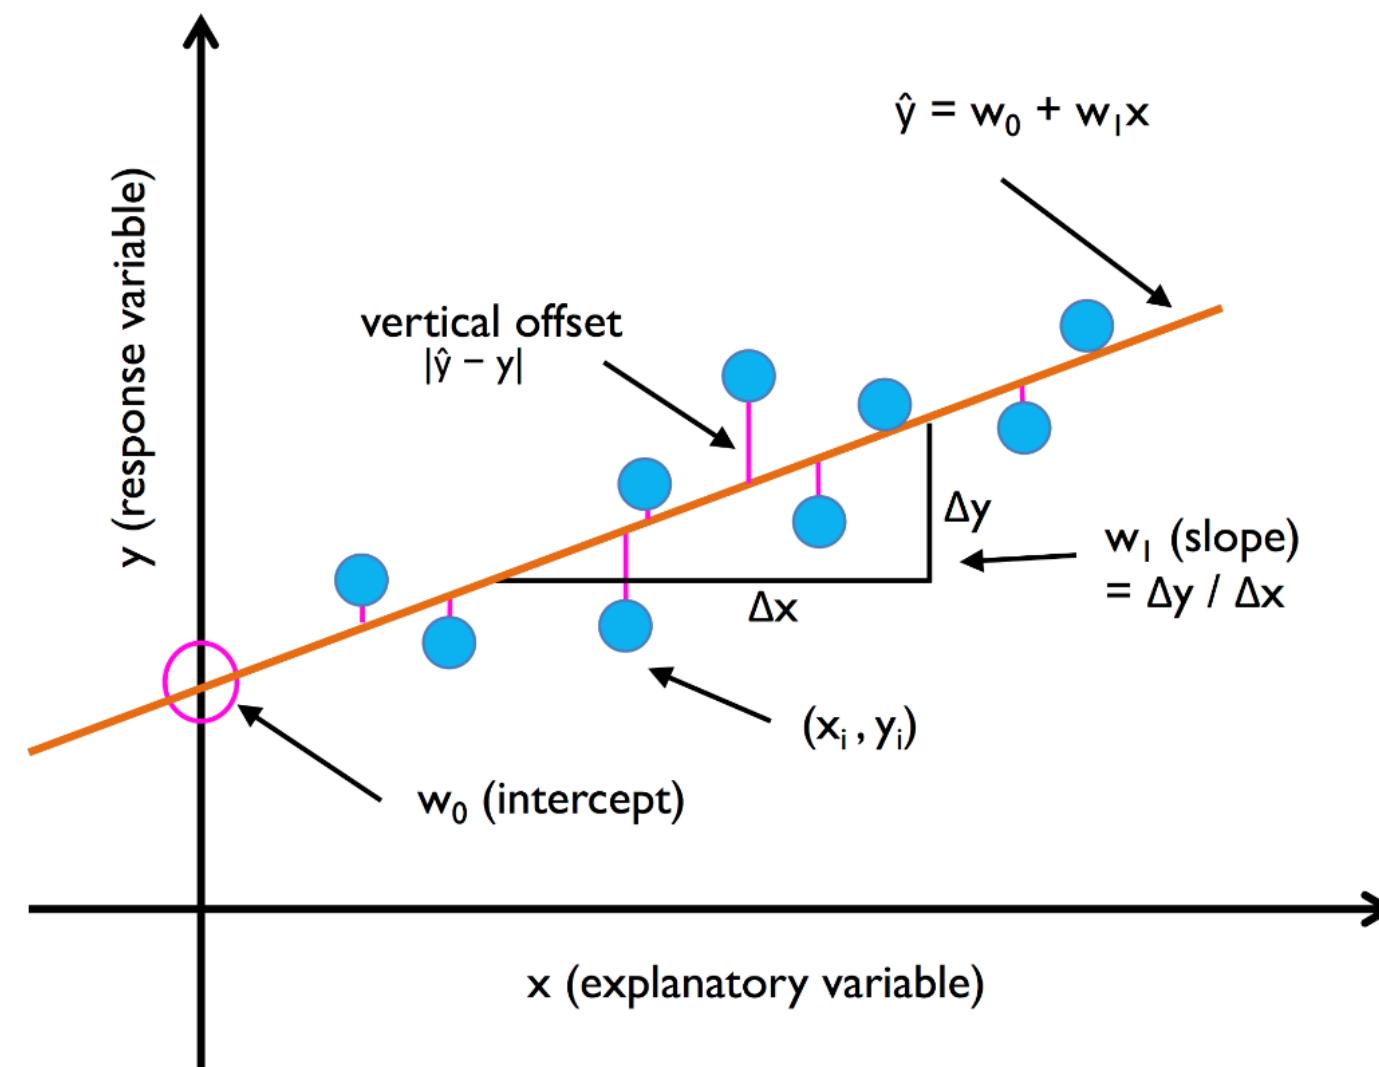
Simple Linear Regression

Simple Linear Regression

$$y_i = w_1 x_i + w_0 + \varepsilon_i$$

- y_i : dependent/endogenous/response target, label (Ex: `alcohol`)
- x_i : independent/exogenous/explanatory feature, attribute (Ex: `proline`)
- w_1 : coefficient, slope
- w_0 : coefficient, bias term, intercept
- ε_i : error, hopefully small, often assumed $\sim \mathcal{N}(0, 1)$
- Want to find values for w_1 and w_0 that best fit the data.
- Find a line as close to our observations as possible

Simple Linear Regression



from PML

Finding w_1 and w_0 with Ordinary Least Squares

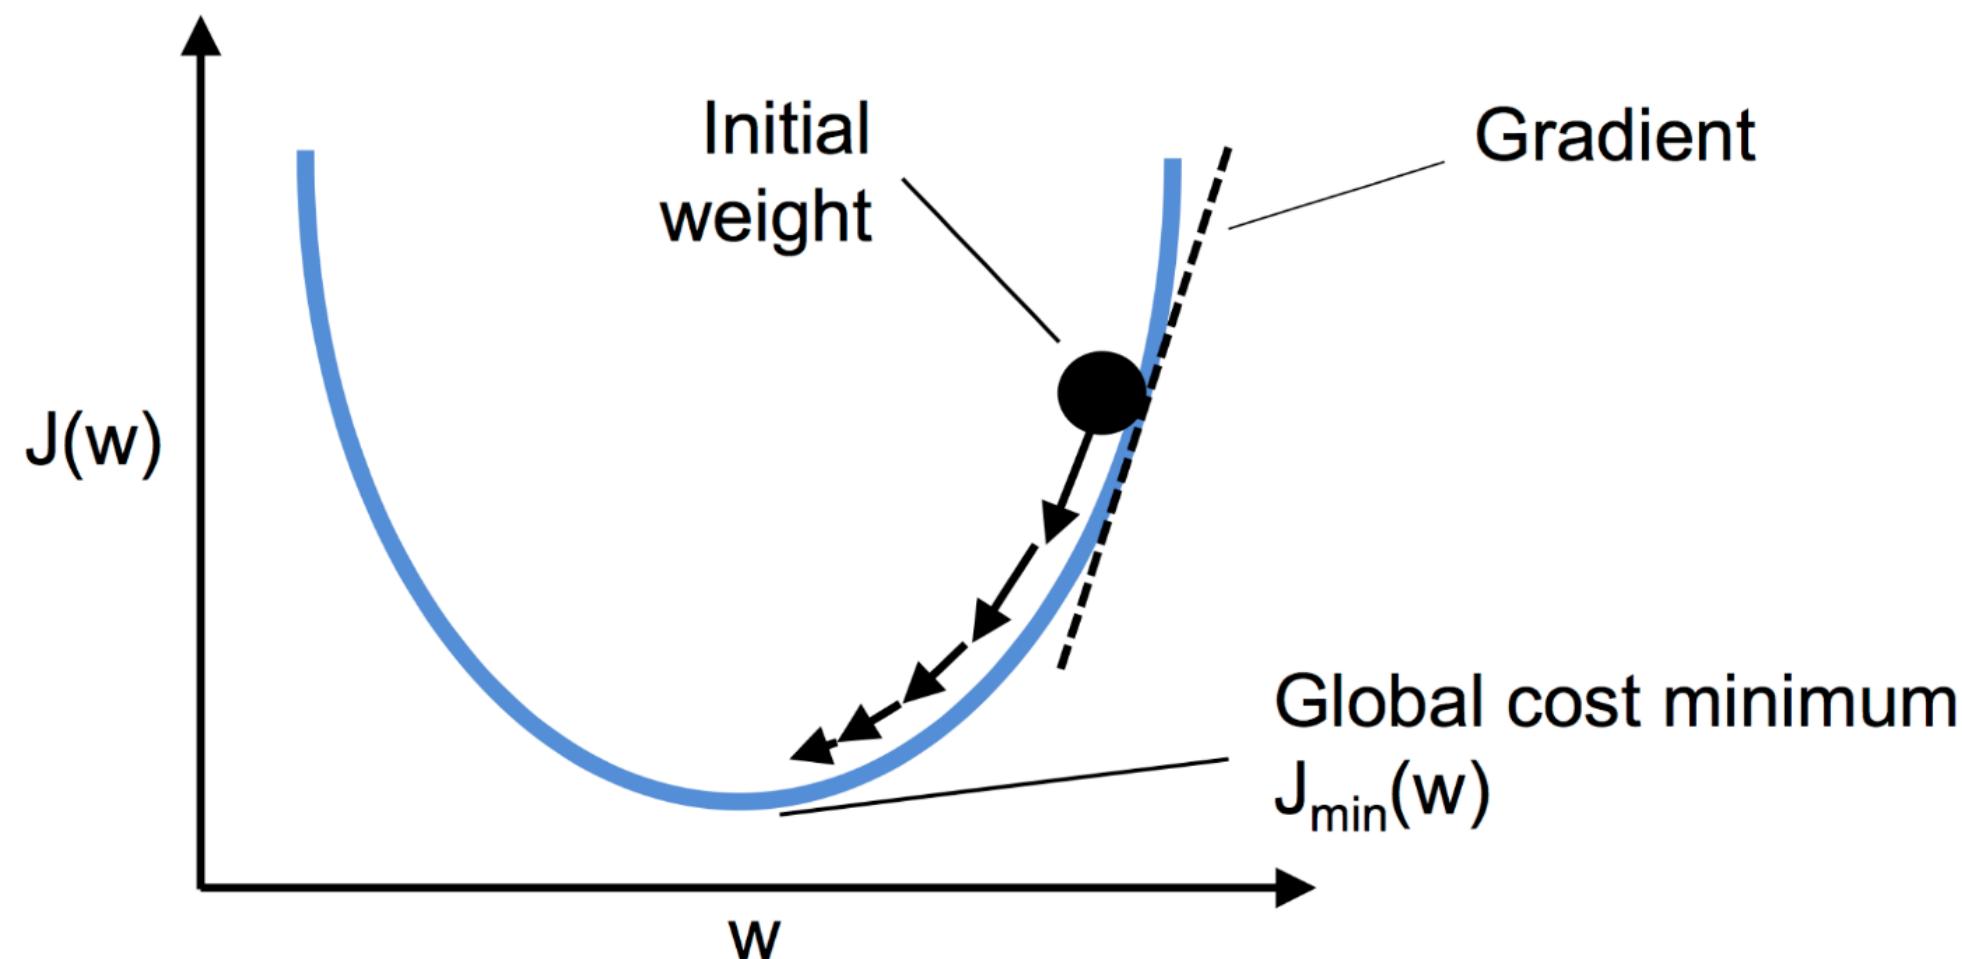
- **prediction:** $\hat{y}_i = f(x_i) = w_1 x_i + w_0$
- **error:** $error(y_i, \hat{y}_i) = y_i - \hat{y}_i$
- **sum of squared errors:** $\sum_{i=1:n} (y_i - \hat{y}_i)^2$
- **least squares:** make the sum of squared errors as small as possible
- **gradient descent:** minimize error by following the gradient wrt w_1, w_0
 - can sometimes be optimized in closed form
 - often done iteratively

Aside: Gradient Descent

- Want to maximize or minimize something (Ex: squared error)
- **Gradient** : direction, vector of partial derivatives
 - can get complicated, often estimated
- **Gradient Descent** : take steps wrt the direction of the gradient
 - **maximize** : in the direction of the gradient
 - **minimize** : in the opposite direction of the gradient
- **Global Maximum/Minimum** : the single best solution
- **Local Maximum/Minimum** : the best solution in the neighborhood

Aside: Gradient Descent Cont.

Simple case with a global min:



Simple Regression Using scikit-learn

Simple Regression Using scikit-learn

In [9]:

```
1 # import the model from sklearn
2 from sklearn.linear_model import LinearRegression
```

Simple Regression Using scikit-learn

```
In [9]: 1 # import the model from sklearn
```

```
2 from sklearn.linear_model import LinearRegression
```

```
In [10]: 1 # instantiate the model and set hyperparameters
```

```
2 lr = LinearRegression(fit_intercept=True)      # by default
```

Simple Regression Using scikit-learn

```
In [9]: 1 # import the model from sklearn
```

```
2 from sklearn.linear_model import LinearRegression
```

```
In [10]: 1 # instantiate the model and set hyperparameters
```

```
2 lr = LinearRegression(fit_intercept=True) # by default
```

```
In [11]: 1 # fit the model
```

```
2 lr.fit(X=df_wine.proline.values.reshape(-1, 1), y=df_wine.alcohol);
```

Simple Regression Using scikit-learn

In [9]:

```
1 # import the model from sklearn  
2 from sklearn.linear_model import LinearRegression
```

In [10]:

```
1 # instantiate the model and set hyperparameters  
2 lr = LinearRegression(fit_intercept=True)      # by default
```

In [11]:

```
1 # fit the model  
2 lr.fit(X=df_wine.proline.values.reshape(-1, 1), y=df_wine.alcohol);
```

In [12]:

```
1 # display learned coefficients (trailing underscore indicates learned values)  
2 print(lr.coef_.round(4))  
3 print(lr.intercept_.round(2))
```

```
[0.0017]  
11.76
```

Simple Regression Using scikit-learn

```
In [9]: 1 # import the model from sklearn
```

```
2 from sklearn.linear_model import LinearRegression
```

```
In [10]: 1 # instantiate the model and set hyperparameters
```

```
2 lr = LinearRegression(fit_intercept=True) # by default
```

```
In [11]: 1 # fit the model
```

```
2 lr.fit(X=df_wine.proline.values.reshape(-1, 1), y=df_wine.alcohol);
```

```
In [12]: 1 # display learned coefficients (trailing underscore indicates learned values)
```

```
2 print(lr.coef_.round(4))
```

```
3 print(lr.intercept_.round(2))
```

```
[0.0017]
```

```
11.76
```

```
In [13]: 1 # predict given new values for proline
```

```
2 X = np.array([1000,2000]).reshape(-1,1)
```

```
3 lr.predict(X).round(2)
```

```
Out[13]: array([13.42, 15.08])
```

Why .reshape(-1,1)?

scikit-learn models expect the input features to be 2 dimensional

Why .reshape(-1,1)?

scikit-learn models expect the input features to be 2 dimensional

```
In [14]: 1 df_wine.proline.values[:5]  
Out[14]: array([1065., 1050., 1185., 1480., 735.])
```

Why .reshape(-1,1)?

scikit-learn models expect the input features to be 2 dimensional

```
In [14]: 1 df_wine.proline.values[:5]
```

```
Out[14]: array([1065., 1050., 1185., 1480., 735.])
```

```
In [15]: 1 df_wine.proline.values.shape
```

```
Out[15]: (178,)
```

Why .reshape(-1,1)?

scikit-learn models expect the input features to be 2 dimensional

```
In [14]: 1 df_wine.proline.values[:5]
```

```
Out[14]: array([1065., 1050., 1185., 1480., 735.])
```

```
In [15]: 1 df_wine.proline.values.shape
```

```
Out[15]: (178,)
```

```
In [16]: 1 df_wine.proline.values.reshape(-1,1).shape # -1 means "infer from the data"
```

```
Out[16]: (178, 1)
```

```
In [17]: 1 # df_wine.proline.values.reshape(-1,1)
```

Why .reshape(-1,1)?

scikit-learn models expect the input features to be 2 dimensional

```
In [14]: 1 df_wine.proline.values[:5]
```

```
Out[14]: array([1065., 1050., 1185., 1480., 735.])
```

```
In [15]: 1 df_wine.proline.values.shape
```

```
Out[15]: (178,)
```

```
In [16]: 1 df_wine.proline.values.reshape(-1,1).shape # -1 means "infer from the data"
```

```
Out[16]: (178, 1)
```

```
In [17]: 1 # df_wine.proline.values.reshape(-1,1)
```

Alternatives:

Why .reshape(-1,1)?

scikit-learn models expect the input features to be 2 dimensional

```
In [14]: 1 df_wine.proline.values[:5]
```

```
Out[14]: array([1065., 1050., 1185., 1480., 735.])
```

```
In [15]: 1 df_wine.proline.values.shape
```

```
Out[15]: (178,)
```

```
In [16]: 1 df_wine.proline.values.reshape(-1,1).shape # -1 means "infer from the data"
```

```
Out[16]: (178, 1)
```

```
In [17]: 1 # df_wine.proline.values.reshape(-1,1)
```

Alternatives:

```
In [18]: 1 df_wine.loc[:,['proline']].shape
```

```
Out[18]: (178, 1)
```

Why .reshape(-1,1)?

scikit-learn models expect the input features to be 2 dimensional

```
In [14]: 1 df_wine.proline.values[:5]
```

```
Out[14]: array([1065., 1050., 1185., 1480., 735.])
```

```
In [15]: 1 df_wine.proline.values.shape
```

```
Out[15]: (178,)
```

```
In [16]: 1 df_wine.proline.values.reshape(-1,1).shape # -1 means "infer from the data"
```

```
Out[16]: (178, 1)
```

```
In [17]: 1 # df_wine.proline.values.reshape(-1,1)
```

Alternatives:

```
In [18]: 1 df_wine.loc[:,['proline']].shape
```

```
Out[18]: (178, 1)
```

```
In [19]: 1 df_wine[['proline']].shape
```

```
Out[19]: (178, 1)
```

Interpreting Coefficients

Interpreting Coefficients

```
In [20]: 1 print(f'w_1 = {lr.coef_[0]:0.3f}, w_0 = {lr.intercept_:0.3f}')
```

```
w_1 = 0.002, w_0 = 11.761
```

Interpreting Coefficients

```
In [20]: 1 print(f'w_1 = {lr.coef_[0]:0.3f}, w_0 = {lr.intercept_:0.3f}')
```

```
w_1 = 0.002, w_0 = 11.761
```

```
In [21]: 1 print(f'alcohol = {lr.coef_[0]:0.3f}*proline + {lr.intercept_:0.3f}')
```

```
alcohol = 0.002*proline + 11.761
```

Interpreting Coefficients

```
In [20]: 1 print(f'w_1 = {lr.coef_[0]:0.3f}, w_0 = {lr.intercept_:0.3f}')
```

```
w_1 = 0.002, w_0 = 11.761
```

```
In [21]: 1 print(f'alcohol = {lr.coef_[0]:0.3f}*proline + {lr.intercept_:0.3f}')
```

```
alcohol = 0.002*proline + 11.761
```

- When proline goes up by 1, alcohol goes up by .002
- When proline is 0, alcohol is 11.761

Plotting The Model

Plotting The Model

In [22]:

```
1 x_predict = [df_wine.proline.min(),df_wine.proline.max()]
2 y_hat = lr.predict(np.array(x_predict).reshape(-1,1))
3
4 fig,ax = plt.subplots(1,1,figsize=(12,8))
5 ax = sns.scatterplot(x=df_wine.proline,y=df_wine.alcohol);
6 ax.plot(x_predict,y_hat);
```



Multiple Linear Regression

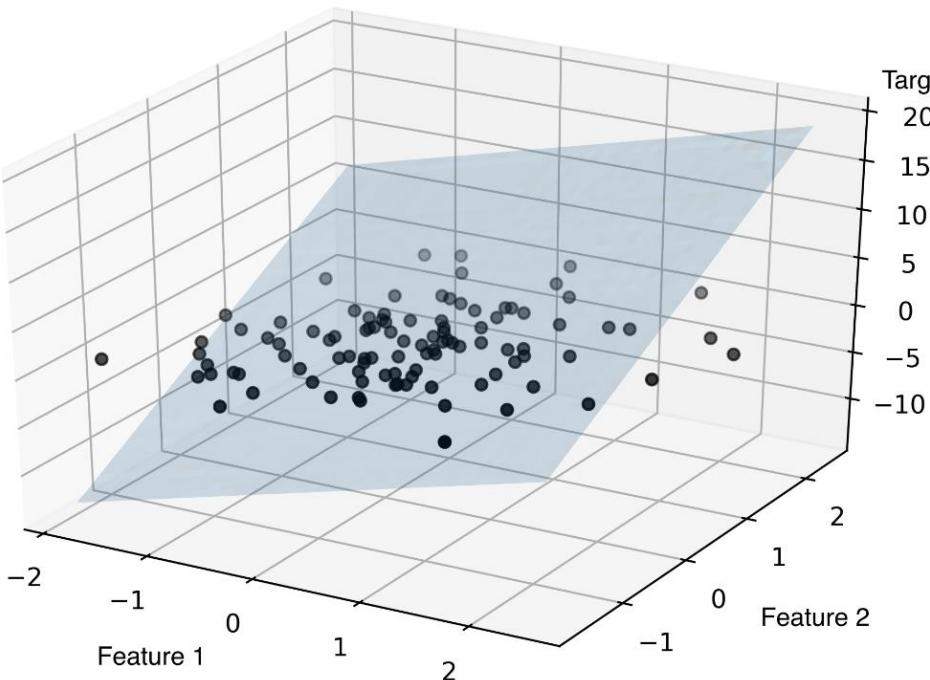
- Including multiple independent variables

$$y_i = w_0 + w_1 x_{i1} + w_2 x_{i2} + \dots + w_m x_{im} + \varepsilon_i$$

Ex:

$$\text{alcohol}_i = w_0 + w_{\text{proline}} \text{proline}_i + w_{\text{hue}} \text{hue}_i$$

Objective: Find a *plane* that falls as close to our points as possible



Multiple Linear Regression in scikit-learn

Multiple Linear Regression in scikit-learn

In [23]:

```
1 mlr = LinearRegression()
2 mlr.fit(df_wine[['proline', 'hue']], y=df_wine.alcohol);
3
4 print(f'{"intercept":10s} : {mlr.intercept_:0.3f}')
5 for (name,coef) in zip(['proline','hue'],mlr.coef_):
6     print(f'{name:10s} : {coef: 0.3f}')
```

```
intercept    : 12.459
proline      :  0.002
hue          : -0.842
```

Multiple Linear Regression in scikit-learn

In [23]:

```
1 mlr = LinearRegression()
2 mlr.fit(df_wine[['proline', 'hue']], y=df_wine.alcohol);
3
4 print(f'{"intercept":10s} : {mlr.intercept_:0.3f}')
5 for (name,coef) in zip(['proline','hue'],mlr.coef_):
6     print(f'{name:10s} : {coef: 0.3f}')
```

```
intercept    : 12.459
proline      :  0.002
hue          : -0.842
```

- If we hold everything else constant, what effect does each variable have?
 - If hue is held constant, a rise of 1 proline -> rise of .002 in alcohol
 - If proline is held constant, a rise of 1 hue -> decrease of .842 in alcohol
- Can add interaction terms to allow both to move at the same time:
 - Ex: hue * proline
 - more complicated to interpret

Multiple Linear Regression in statsmodels

Multiple Linear Regression in statsmodels

In [24]:

```
1 import statsmodels.api as sm
2
3 X = df_wine[['proline', 'hue']].copy()
4 X = sm.add_constant(X) # or X['const'] = 1
5 y = df_wine.alcohol
6 sm_mlr = sm.OLS(y,X).fit() # Note: X,y passed as parameters to object, not fit
7 sm_mlr.summary()
```

Out[24]: OLS Regression Results

Dep. Variable:	alcohol	R-squared:	0.467
Model:	OLS	Adj. R-squared:	0.461
Method:	Least Squares	F-statistic:	76.79
Date:	Tue, 07 Oct 2025	Prob (F-statistic):	1.15e-24
Time:	20:34:38	Log-Likelihood:	-158.89
No. Observations:	178	AIC:	323.8
Df Residuals:	175	BIC:	333.3
Df Model:	2		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	12.4593	0.203	61.347	0.000	12.058	12.860
proline	0.0018	0.000	12.325	0.000	0.002	0.002
hue	-0.8418	0.202	-4.175	0.000	-1.240	-0.444

Omnibus:	0.751	Durbin-Watson:	1.734
Prob(Omnibus):	0.687	Jarque-Bera (JB):	0.606
Skew:	0.142	Prob(JB):	0.739
Kurtosis:	3.028	Cond. No.	4.96e+03

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Dealing With the Intercept/Bias

- Two ways of keeping track of the bias term

Dealing With the Intercept/Bias

- Two ways of keeping track of the bias term

1. Keep it as a separate parameter (sklearn) :

- $y = w_0 + w_1x_1 + w_2x_2 + \dots + w_mx_m$
- $y = w_0 + \sum_{i=1}^m w_i x_i$

Dealing With the Intercept/Bias

- Two ways of keeping track of the bias term

1. Keep it as a separate parameter (sklearn) :

- $y = w_0 + w_1x_1 + w_2x_2 + \dots + w_mx_m$
- $y = w_0 + \sum_{i=1}^m w_i x_i$

2. Append a constant of $x_0 = 1$ so x and w are the same length (statsmodels) :

- $y = w_0x_0 + w_1x_1 + w_2x_2 + \dots + w_mx_m$
- $y = \sum_{i=0}^m w_i x_i$

Dealing With the Intercept/Bias

- Two ways of keeping track of the bias term

1. Keep it as a separate parameter (sklearn) :

- $y = w_0 + w_1x_1 + w_2x_2 + \dots + w_mx_m$
- $y = w_0 + \sum_{i=1}^m w_i x_i$

2. Append a constant of $x_0 = 1$ so x and w are the same length (statsmodels) :

- $y = w_0x_0 + w_1x_1 + w_2x_2 + \dots + w_mx_m$
- $y = \sum_{i=0}^m w_i x_i$

```
In [25]: 1 | x.head(3)
```

```
Out[25]:
```

	const	proline	hue
0	1.0	1065.0	1.04
1	1.0	1050.0	1.05
2	1.0	1185.0	1.03

Standardizing/Normalizing Features for Interpretation

Standardizing/Normalizing Features for Interpretation

```
In [26]: 1 for (name,coef) in zip(['proline','hue'],mlr.coef_):  
2     print(f'{name:10s} : {coef: 0.3f}')
```

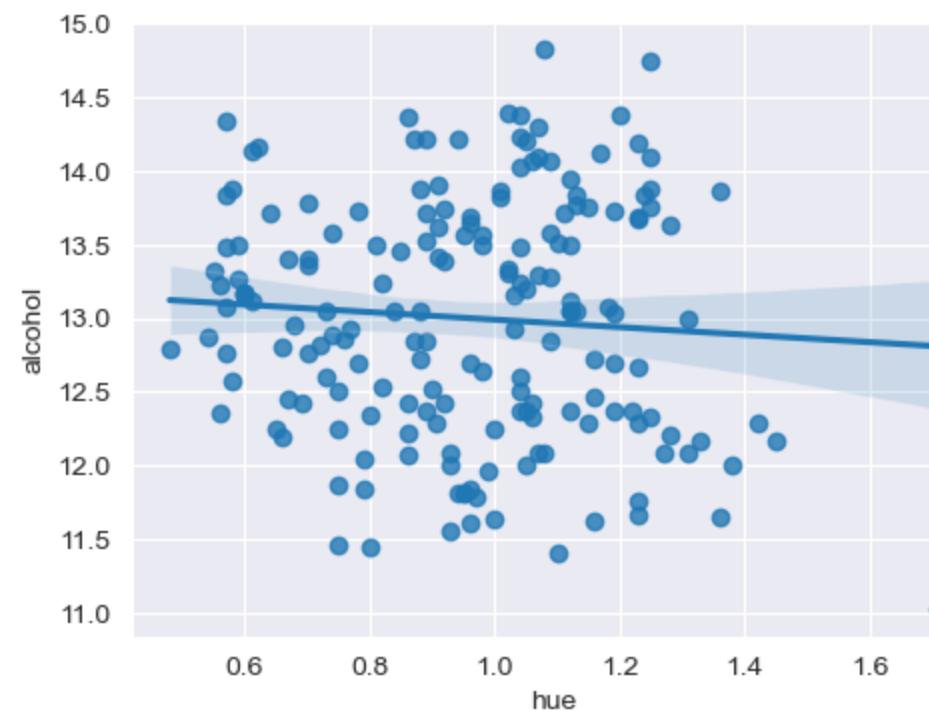
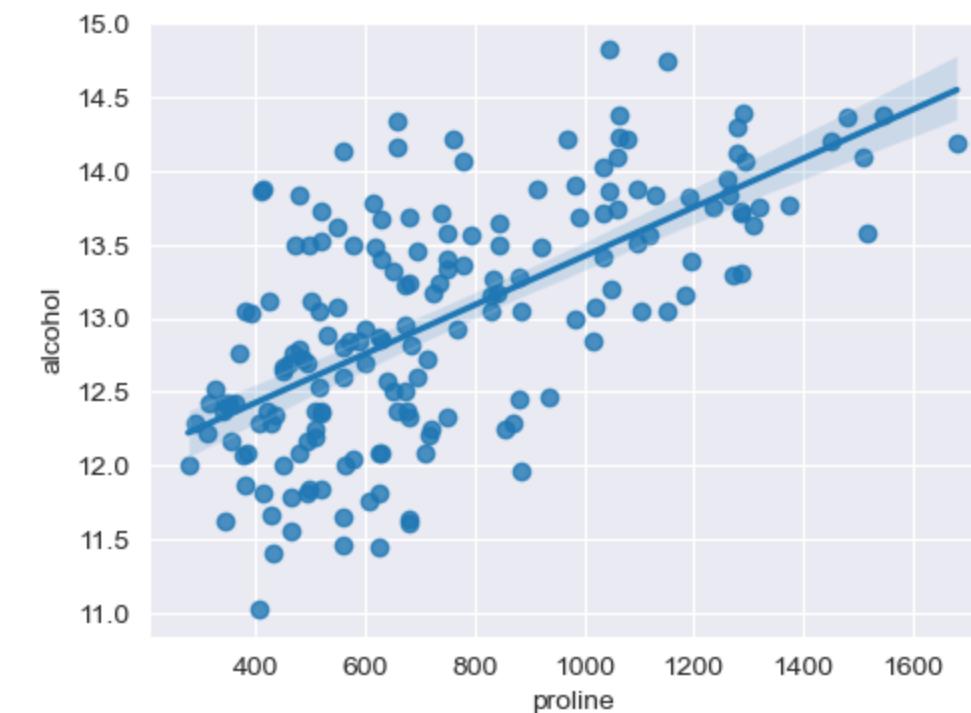
```
proline      :  0.002  
hue          : -0.842
```

Standardizing/Normalizing Features for Interpretation

```
In [26]: 1 for (name,coef) in zip(['proline','hue'],mlr.coef_):  
2     print(f'{name:10s} : {coef: 0.3f}')
```

```
proline      :  0.002  
hue          : -0.842
```

```
In [27]: 1 fig,ax = plt.subplots(1,2,figsize=(12,4))  
2 sns.regplot(x='proline',y='alcohol',data=df_wine,ax=ax[0])  
3 sns.regplot(x='hue',y='alcohol',data=df_wine,ax=ax[1]);
```

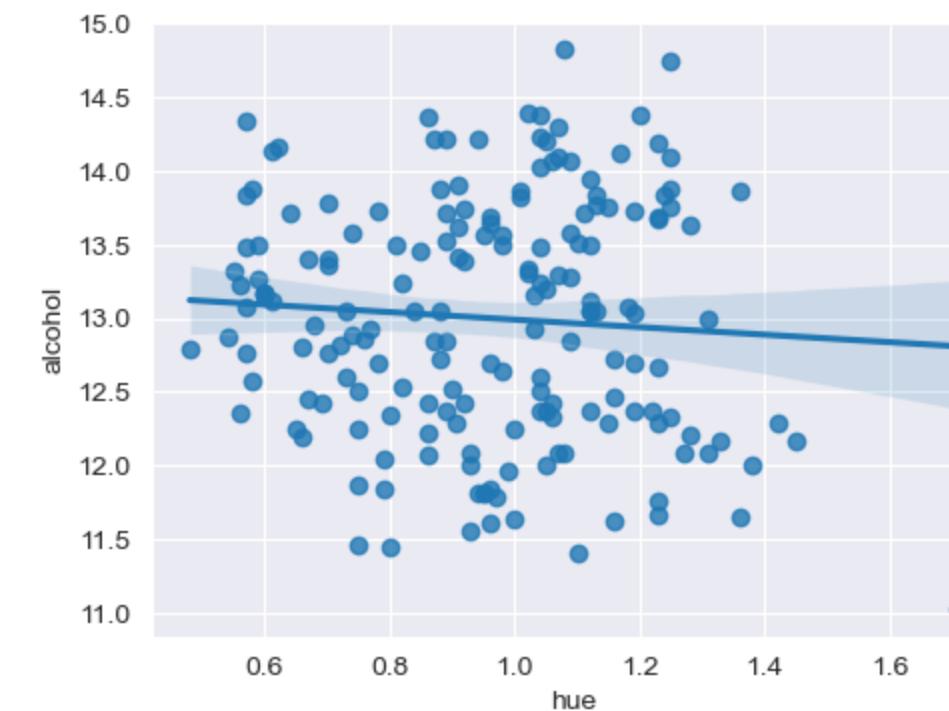
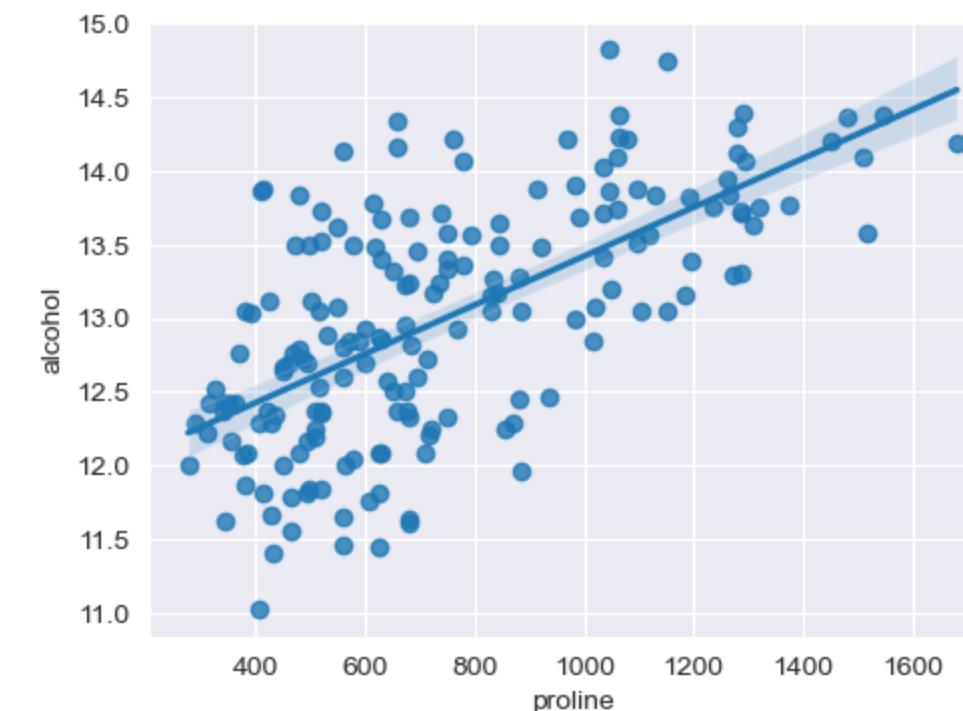


Standardizing/Normalizing Features for Interpretation

```
In [26]: 1 for (name,coef) in zip(['proline','hue'],mlr.coef_):  
2     print(f'{name:10s} : {coef: 0.3f}')
```

```
proline      :  0.002  
hue          : -0.842
```

```
In [27]: 1 fig,ax = plt.subplots(1,2,figsize=(12,4))  
2 sns.regplot(x='proline',y='alcohol',data=df_wine,ax=ax[0])  
3 sns.regplot(x='hue',y='alcohol',data=df_wine,ax=ax[1]);
```



What would the coefficients look like if the features were on the same scale?

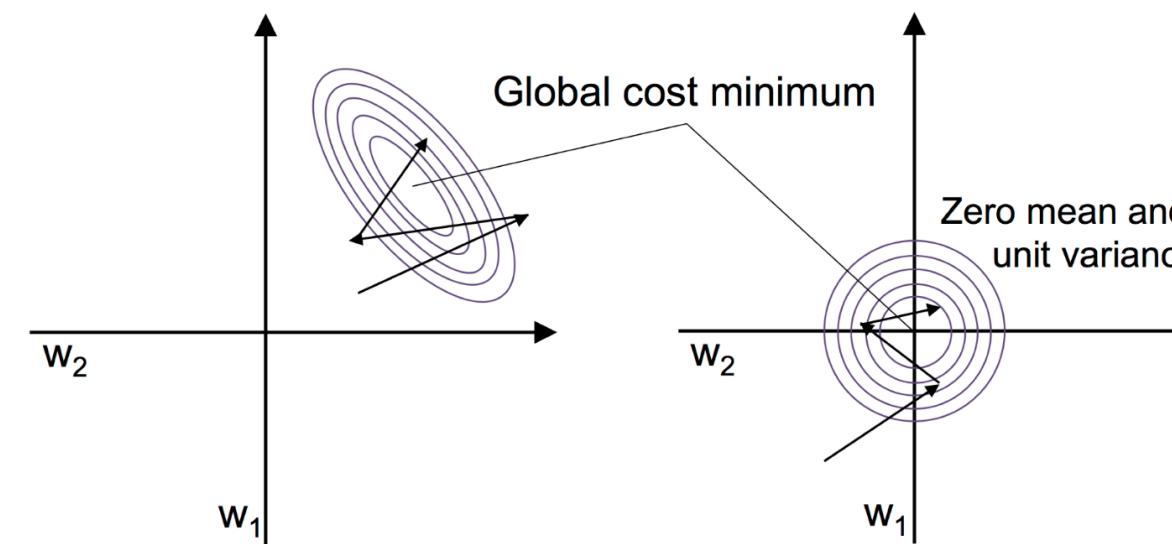
Standardizing/Normalizing Features for Gradient Descent

Standardizing/Normalizing Features for Gradient Descent

$$z = \frac{x - \bar{x}}{s}$$

Standardizing/Normalizing Features for Gradient Descent

$$z = \frac{x - \bar{x}}{s}$$



From PML

Multiple Linear Regression with Standardization/Normalization

Multiple Linear Regression with Standardization/Normalization

- `DataFrame.apply()`: apply a function across rows (`axis=0`) or columns (`axis=1`)

Multiple Linear Regression with Standardization/Normalization

- `DataFrame.apply()`: apply a function across rows (`axis=0`) or columns (`axis=1`)

```
In [28]: 1 X_zscore = df_wine[['proline', 'hue']].apply(lambda x: (x-x.mean())/x.std(), axis=0) # or use StandardScaler
2
3 mlr_n = LinearRegression()
4 mlr_n.fit(X_zscore, df_wine.alcohol)
5 for (name,coef) in zip(X_zscore.columns,mlr_n.coef_):
6     print(f'{name:10s} : {coef: 0.3f}')
```

```
proline      :  0.568
hue          : -0.192
```

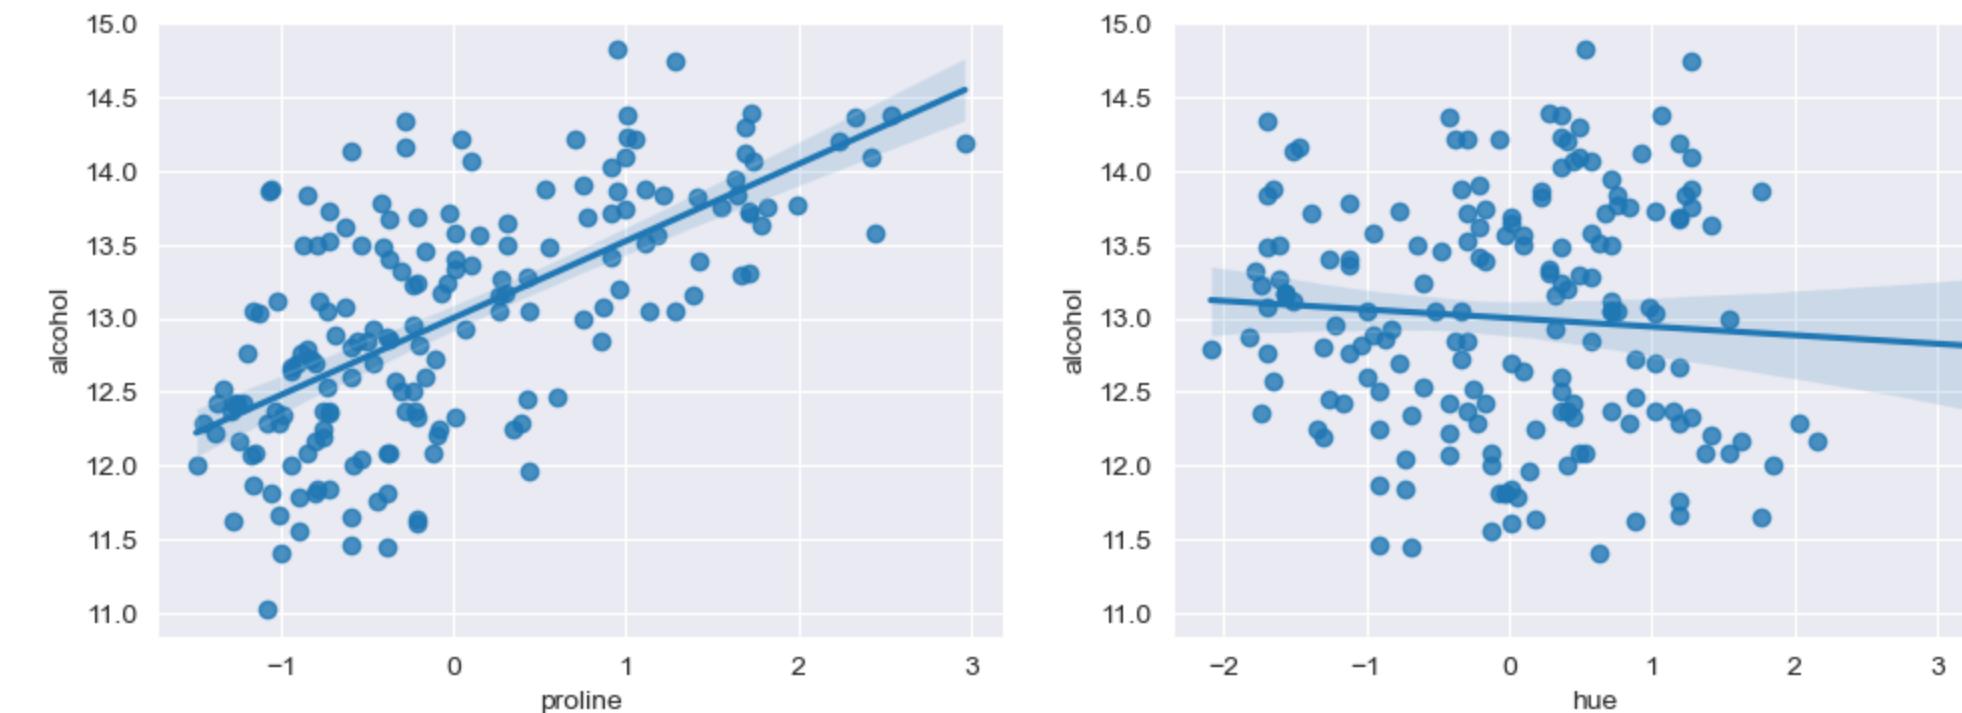
Multiple Linear Regression with Standardization/Normalization

- `DataFrame.apply()`: apply a function across rows (`axis=0`) or columns (`axis=1`)

```
In [28]: 1 X_zscore = df_wine[['proline', 'hue']].apply(lambda x: (x-x.mean())/x.std(), axis=0) # or use StandardScaler  
2  
3 mlr_n = LinearRegression()  
4 mlr_n.fit(X_zscore, df_wine.alcohol)  
5 for (name,coef) in zip(X_zscore.columns,mlr_n.coef_):  
6     print(f'{name}: {10}s : {coef: 0.3f}')
```

proline : 0.568
hue : -0.192

```
In [29]: 1 fig,ax = plt.subplots(1,2,figsize=(12,4))  
2 sns.regplot(x=X_zscore.proline,y=df_wine.alcohol,ax=ax[0]);  
3 sns.regplot(x=X_zscore.hue,y=df_wine.alcohol,ax=ax[1]);
```



Aside: apply(axis=), rows vs columns

	proline	hue
0	1065.0	1.04
1	1050.0	1.05
2	1185.0	1.03
3	1480.0	0.86
4	735.0	1.04

axis=0

or

axis='rows'

	proline	hue
0	1065.0	1.04
1	1050.0	1.05
2	1185.0	1.03
3	1480.0	0.86
4	735.0	1.04

axis=1

or

axis='columns'

Colinarity

- MLR assumes features are **linearly independent**
 - Can't rewrite one column as a weighted sum of the others
 - Ex: at a restaurant, "proportion of people ordering pasta" likely not related to "table size"?
 - **linearly dependent** Ex: "number of entrees ordered" and "table size"
- Issue: With Linearly Dependent features, the model won't know how to estimate w
 - If we add to one w_i and subtract from another, there will be no change in error
- Try to remove obvious colinearity
 - can use correlation and linear regression to detect
 - Important to consider when constructing categorical features (feature engineering)

Colinarity

- MLR assumes features are **linearly independent**
 - Can't rewrite one column as a weighted sum of the others
 - Ex: at a restaurant, "proportion of people ordering pasta" likely not related to "table size"?
 - **linearly dependent** Ex: "number of entrees ordered" and "table size"
- Issue: With Linearly Dependent features, the model won't know how to estimate w
 - If we add to one w_i and subtract from another, there will be no change in error
- Try to remove obvious colinearity
 - can use correlation and linear regression to detect
 - Important to consider when constructing categorical features (feature engineering)

```
In [30]: 1 df_wine.corr().style.background_gradient().format(' {:.2f} ')
```

Out[30]:

	alcohol	ash	hue	proline	class
alcohol	1.00	0.21	-0.07	0.64	-0.33
ash	0.21	1.00	-0.07	0.22	-0.05
hue	-0.07	-0.07	1.00	0.24	-0.62
proline	0.64	0.22	0.24	1.00	-0.63
class	-0.33	-0.05	-0.62	-0.63	1.00

Aside: Interpretation Vs. Prediction

- Interpretation: Explain how observed features relate to observed target
- Prediction: Given a new observation, can we generate a prediction
- Often asked to do one or the other, be clear which is most important
- In prediction, may not worry about interpreting the model!
- In general, there is increased attention on interpretability

Questions re Regression with Linear Models?

Classification

- So far: **Regression** -> predict a numeric value
- Next: **Classification** -> predict a discrete category/class
- **Binary classification** : two categories
 - pos/neg, cat/dog, win/lose
- **Multiclass classification** : more than two categories/classes
 - red/green/blue, flower type, integer 0-10
- **Multilabel classification** : can assign more than one label to an instance
 - paper topics, entities in image

Wine as Binary Classification

Wine as Binary Classification

```
In [31]: 1 df_wine['class'].value_counts() # originally multiclass
```

```
Out[31]: 1    71
          0    59
          2    48
Name: class, dtype: int64
```

Wine as Binary Classification

```
In [31]: 1 df_wine['class'].value_counts() # originally multiclass
```

```
Out[31]: 1    71
0    59
2    48
Name: class, dtype: int64
```

```
In [32]: 1 # only keep classes 0 and 1
2 df_wine_2class = df_wine[df_wine['class'] < 2]
3
4 # rename 'class' as 'target', since class is a reserved python word
5 df_wine_2class = df_wine_2class.rename({'class':'target'},axis=1)
6
7 df_wine_2class.target.value_counts()
```

```
Out[32]: 1    71
0    59
Name: target, dtype: int64
```

Classifying Wine with a Linear Model

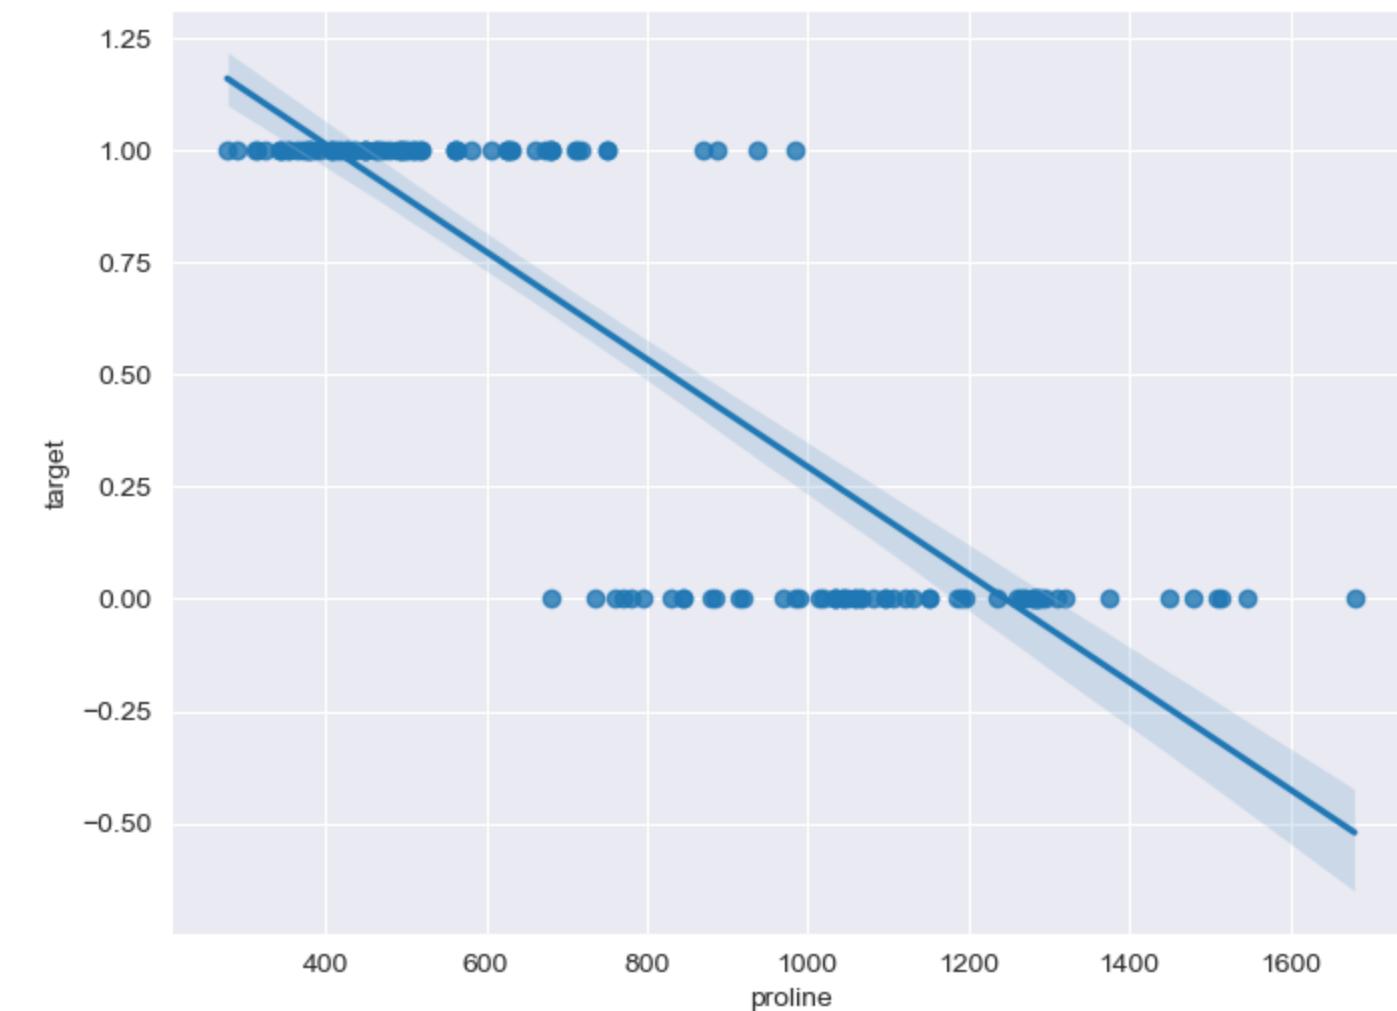
Classifying Wine with a Linear Model

- Can't use our linear regression model directly

Classifying Wine with a Linear Model

- Can't use our linear regression model directly

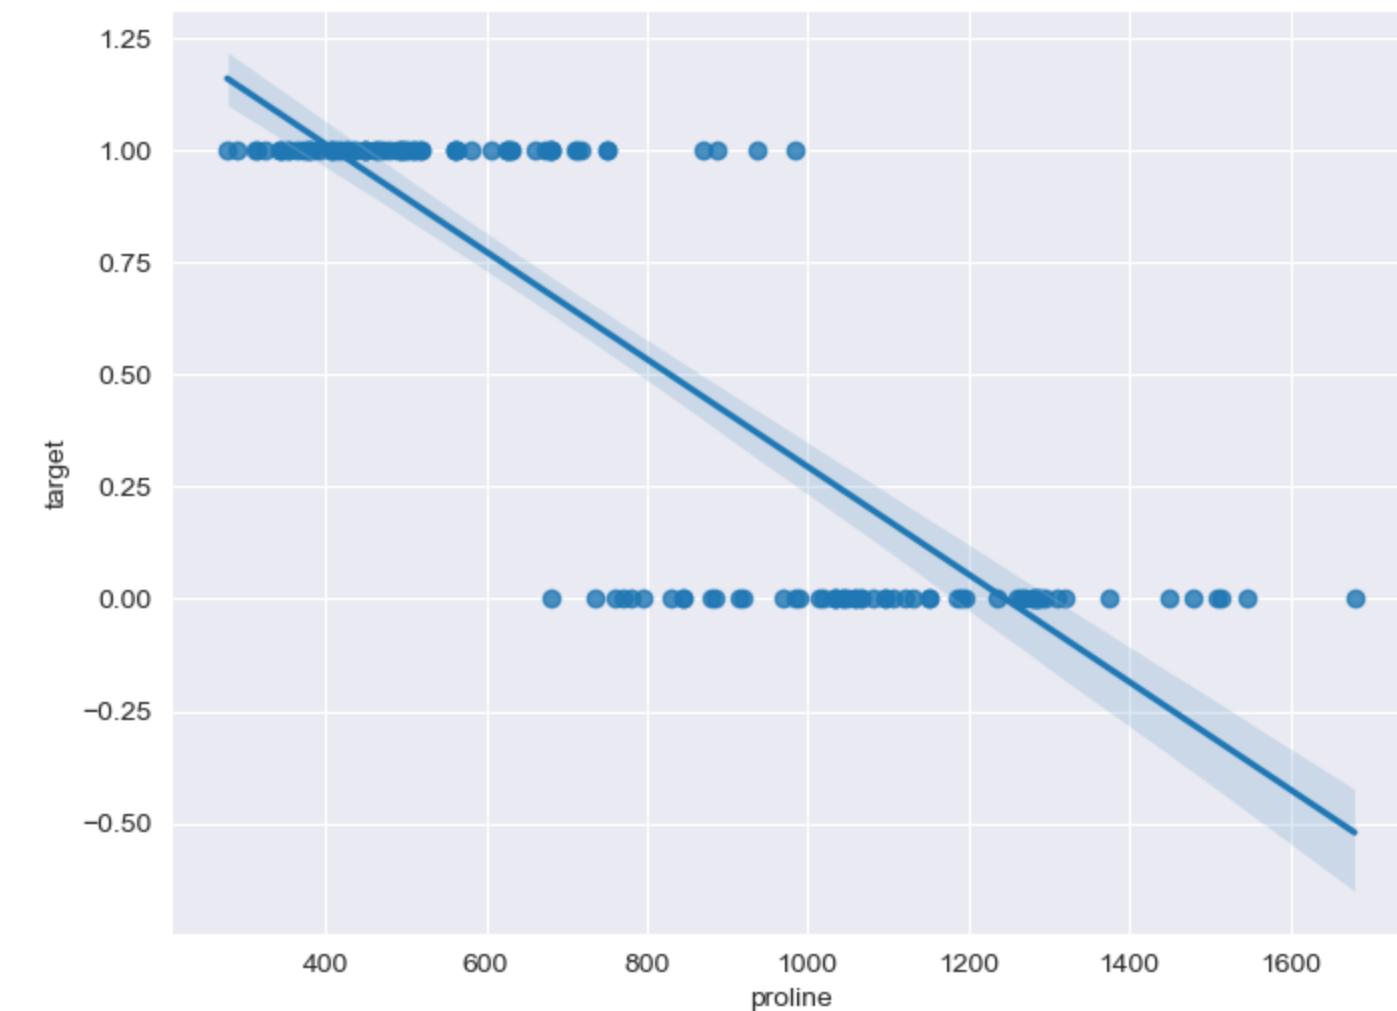
```
In [33]: 1 fig,ax = plt.subplots(1,1,figsize=(8,6))
2 sns.regplot(x=df_wine_2class.proline,y=df_wine_2class.target);
3 #ax.plot([400,800,800,1600],[1,1,0,0],c='r'); # can we add a threshold?
```



Classifying Wine with a Linear Model

- Can't use our linear regression model directly

```
In [33]: 1 fig,ax = plt.subplots(1,1,figsize=(8,6))
2 sns.regplot(x=df_wine_2class.proline,y=df_wine_2class.target);
3 #ax.plot([400,800,800,1600],[1,1,0,0],c='r'); # can we add a threshold?
```



- Want something with that looks like a threshold
- Would like a prediction between 0 and 1

Logistic Regression

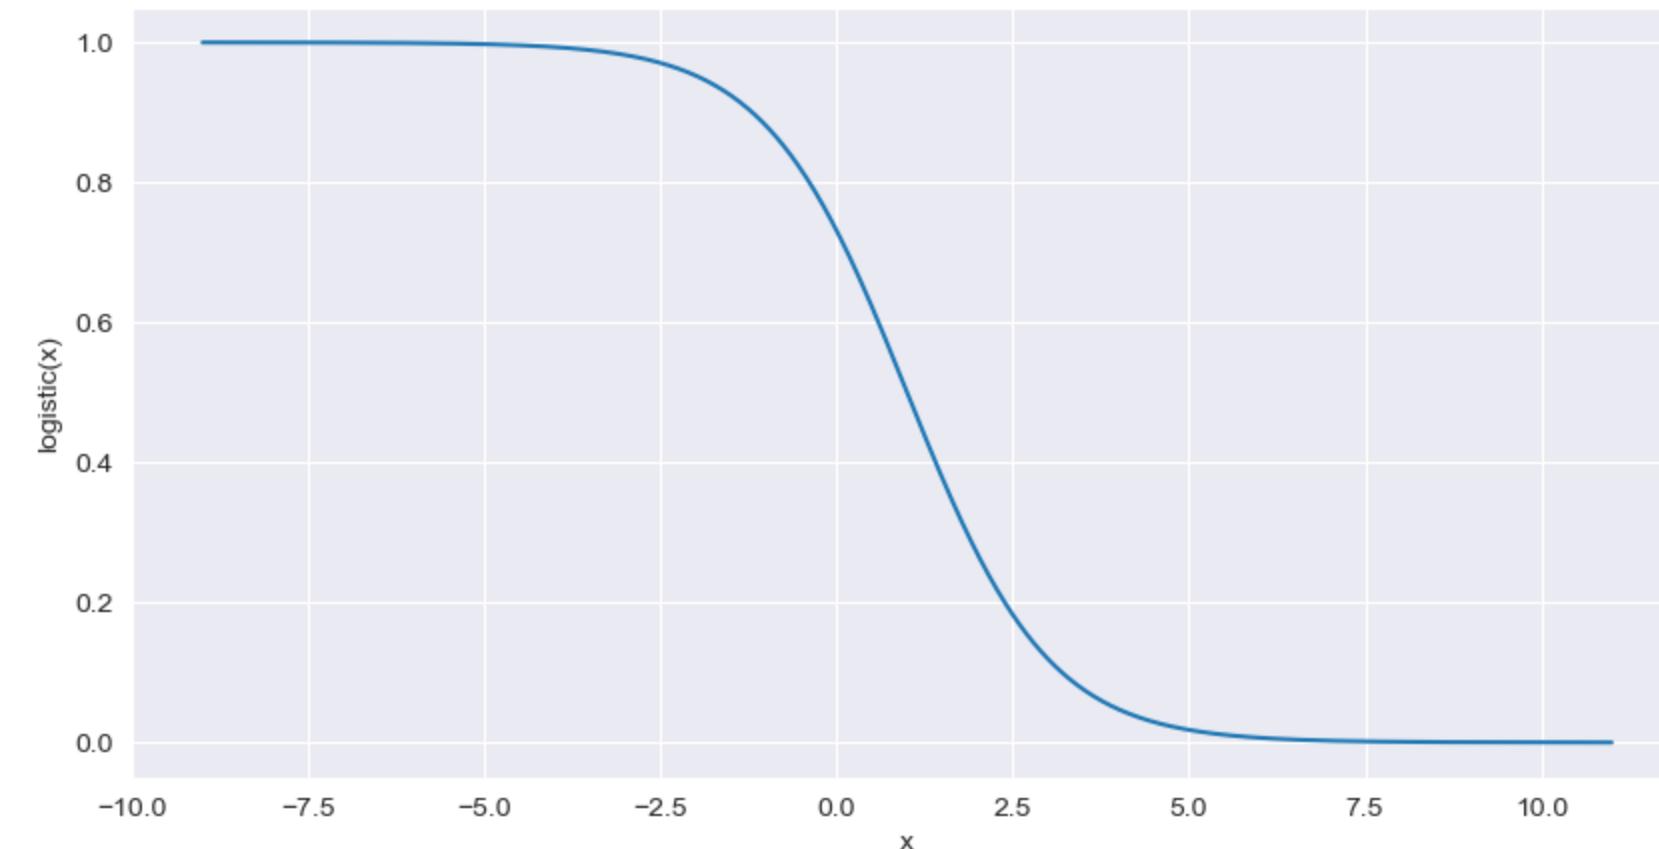
Logistic Regression

$$\text{logistic}(x) = \frac{1}{1+e^{(-x)}}$$

Logistic Regression

$$\text{logistic}(x) = \frac{1}{1+e^{(-x)}}$$

```
In [34]: 1 def logistic(x,w1=1,w0=0):
2     return 1 / (1+np.exp(-(w0+w1*x)))
3
4 x = np.linspace(-10,10,1000) # generate 1000 numbers evenly spaced between -10 and 10
5 fig,ax = plt.subplots(1,1,figsize=(10,5))
6 ax.plot(1-x,logistic(x));
7 ax.set_xlabel('x');ax.set_ylabel('logistic(x)');
8
```



Logistic Regression with sklearn

Logistic Regression with sklearn

- Our problem becomes: $P(y_i = 1|x_i) = \text{logistic}(w_0 + w_1 x_i) + \varepsilon_i$

Logistic Regression with sklearn

- Our problem becomes: $P(y_i = 1|x_i) = \text{logistic}(w_0 + w_1 x_i) + \varepsilon_i$

```
In [35]: 1 from sklearn.linear_model import LogisticRegression
2
3 X_proline = df_wine_2class.proline.values.reshape(-1,1)
4 y_cl = df_wine_2class.target
5
6 logr = LogisticRegression(fit_intercept=True).fit(X_proline,y_cl)
7 print(f'w_0 = {logr.intercept_[0]:0.2f}')
8 print(f'w_1 = {logr.coef_[0][0]:0.2f}')


w_0 = 11.97
w_1 = -0.01
```

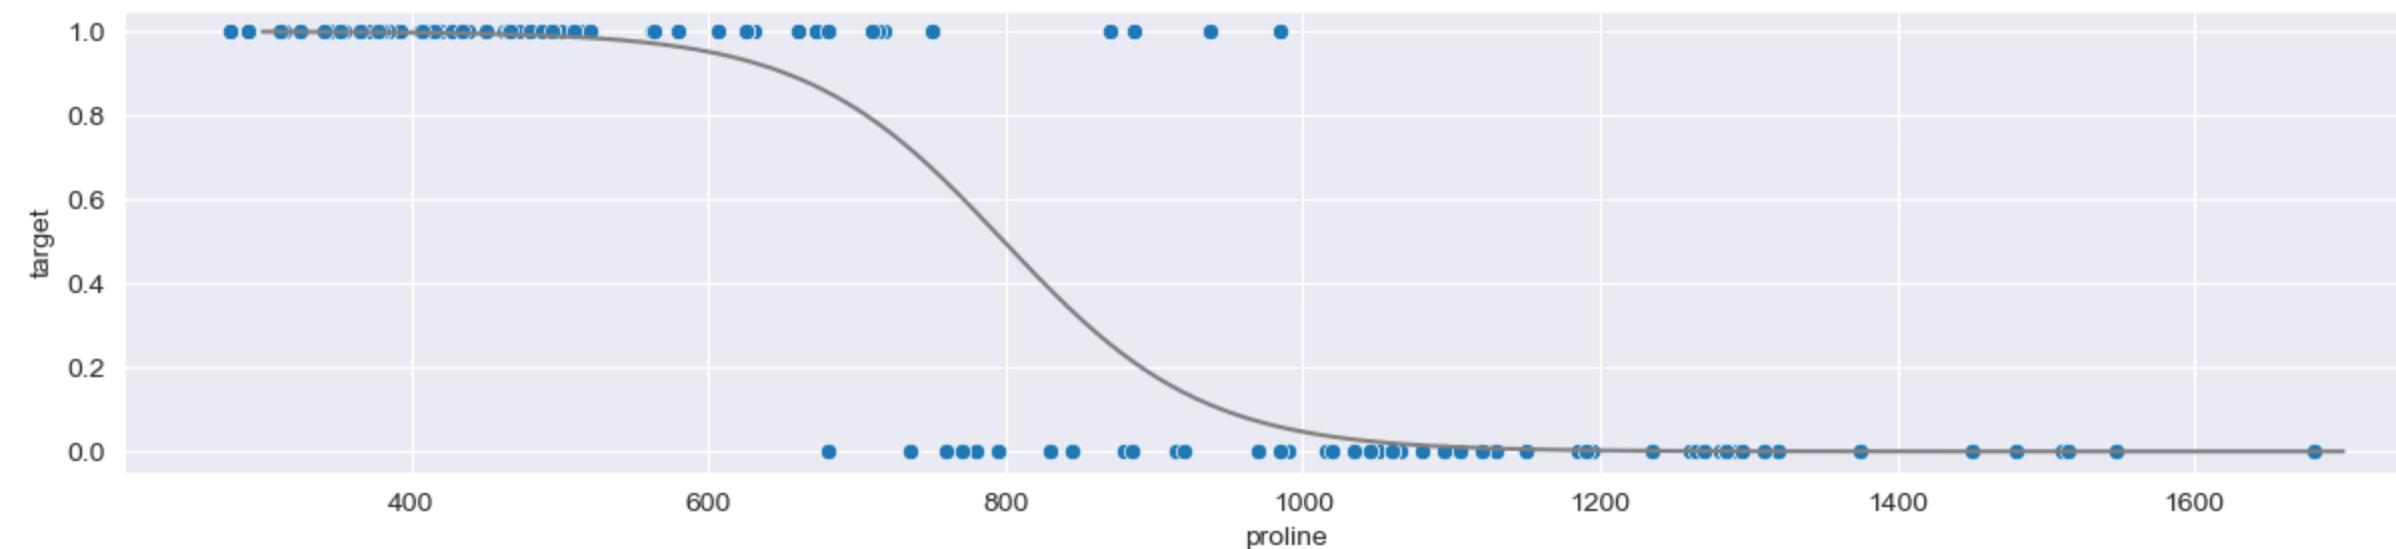
Logistic Regression with sklearn

- Our problem becomes: $P(y_i = 1|x_i) = \text{logistic}(w_0 + w_1 x_i) + \varepsilon_i$

```
In [35]: 1 from sklearn.linear_model import LogisticRegression  
2  
3 X_proline = df_wine_2class.proline.values.reshape(-1,1)  
4 y_cl = df_wine_2class.target  
5  
6 logr = LogisticRegression(fit_intercept=True).fit(X_proline,y_cl)  
7 print(f'w_0 = {logr.intercept_[0]:0.2f}')  
8 print(f'w_1 = {logr.coef_[0][0]:0.2f}')
```

```
w_0 = 11.97  
w_1 = -0.01
```

```
In [36]: 1 fig,ax = plt.subplots(1,1,figsize=(15,3))  
2 x = np.linspace(300,1700,1000)  
3 logistic_x = logistic(x,logr.coef_[0],logr.intercept_)  
4 ax.plot(x,logistic_x,c='gray');  
5 sns.scatterplot(x=df_wine_2class.proline,y=df_wine_2class.target, ax=ax);
```



Adding the Threshold

Adding the Threshold

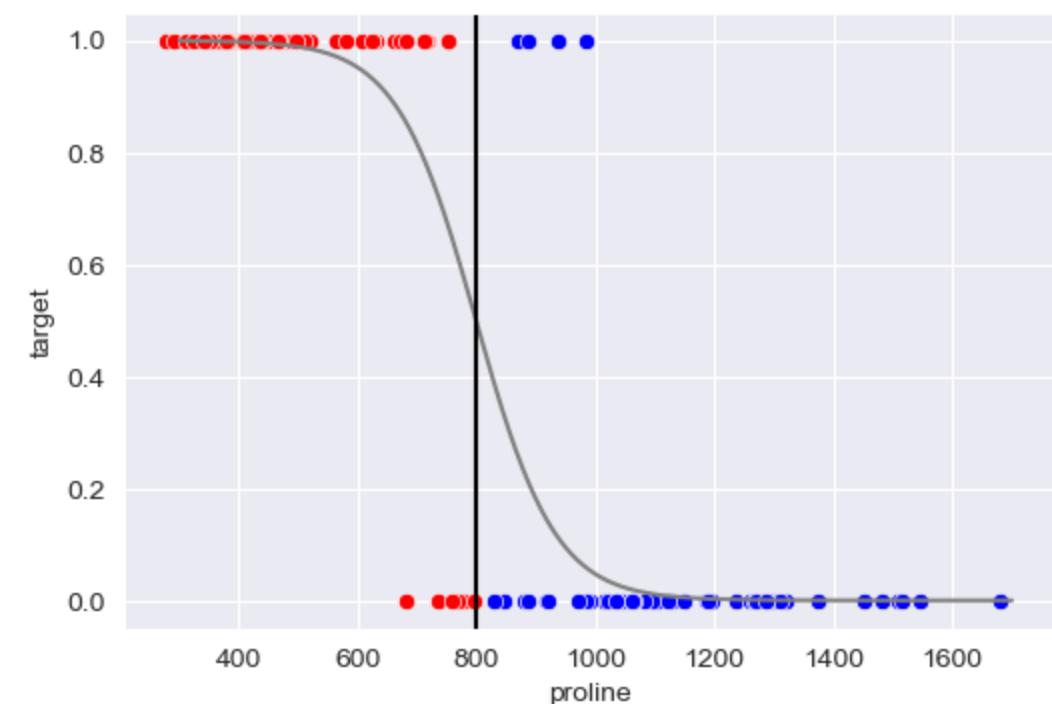
- Can treat the output of the logistic function as $P(y = 1|x)$
- Threshold at .5 (50%) to get class prediction

Adding the Threshold

- Can treat the output of the logistic function as $P(y = 1|x)$
- Threshold at .5 (50%) to get class prediction

```
In [37]: 1 threshold = x[np.argmin(np.abs(logistic_x - .5))]

2
3 predicted_0 = df_wine_2class[df_wine_2class.proline <= threshold]
4 predicted_1 = df_wine_2class[df_wine_2class.proline > threshold]
5
6 fig,ax = plt.subplots(1,1,figsize=(6,4))
7 sns.scatterplot(x='proline',y='target', data=predicted_0, color='r',ax=ax);
8 sns.scatterplot(x='proline',y='target', data=predicted_1, color='b',ax=ax);
9 ax.plot(x,logistic_x,c='gray');
10 ax.axvline(threshold,c='k');
```

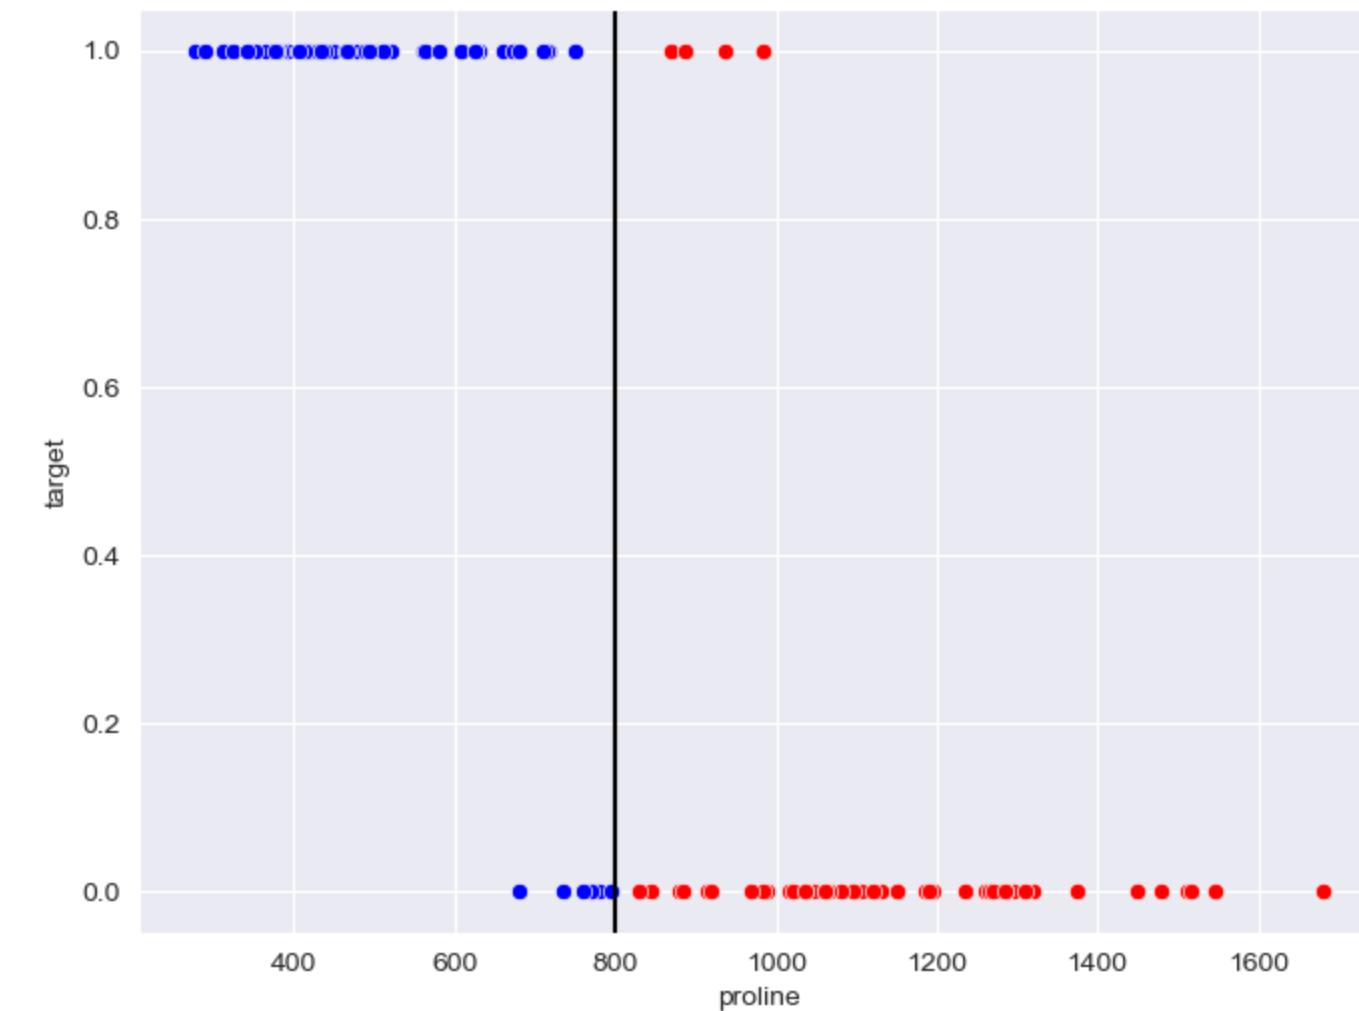


Getting Predictions from sklearn

Getting Predictions from sklearn

In [38]:

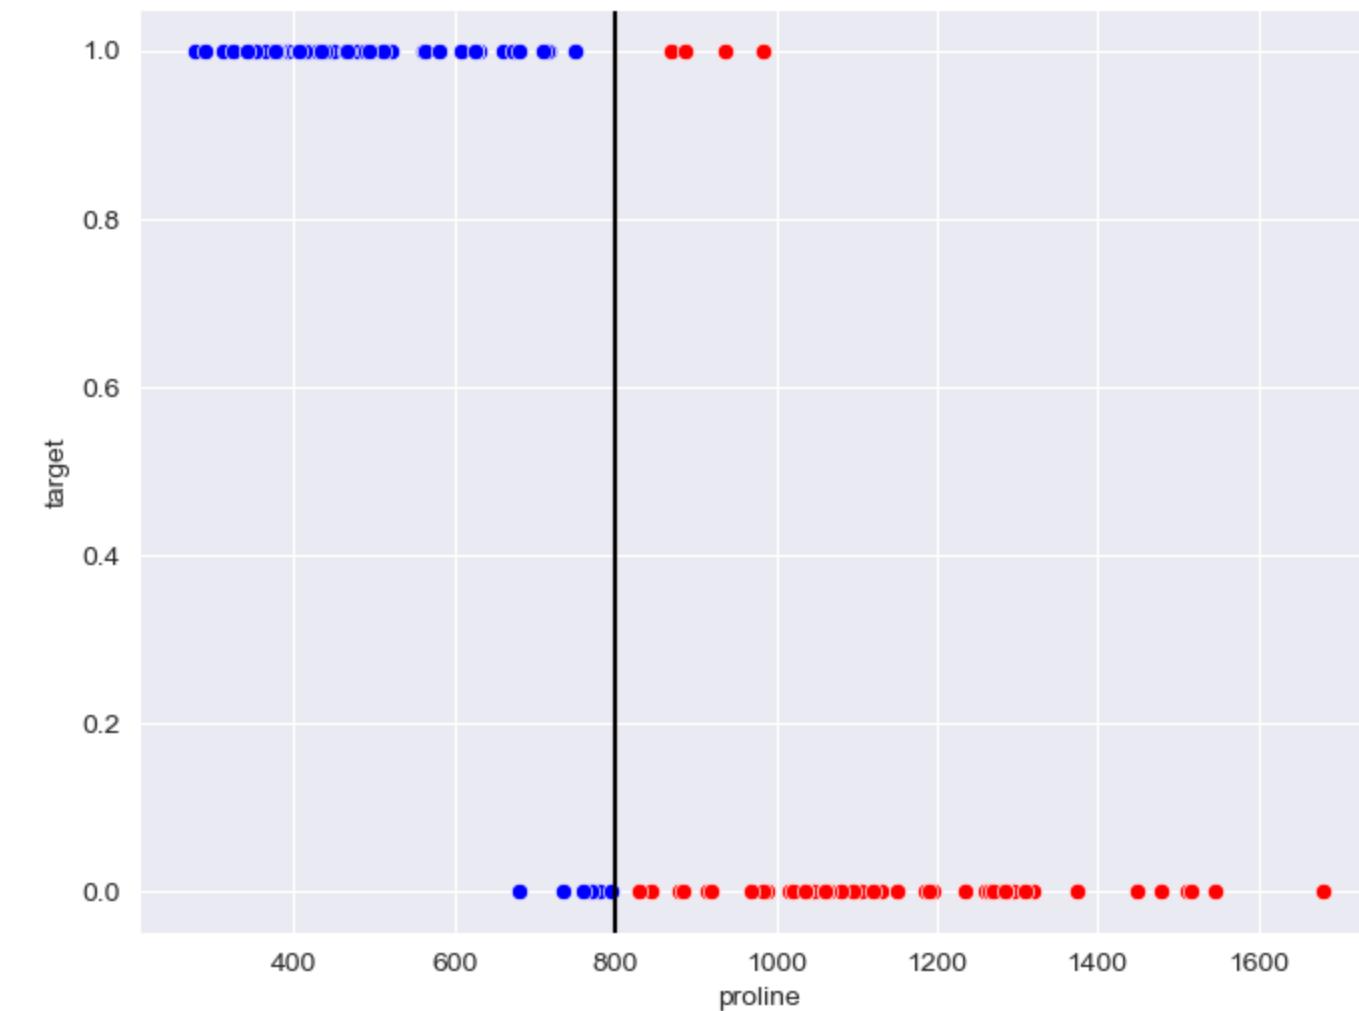
```
1 yhat_cl = logr.predict(X_proline)
2
3 predicted_0 = df_wine_2class[yhat_cl==0]
4 predicted_1 = df_wine_2class[yhat_cl==1]
5
6 fig,ax = plt.subplots(1,1,figsize=(8,6))
7 sns.scatterplot(x='proline',y='target', data=predicted_0, color='r',ax=ax);
8 sns.scatterplot(x='proline',y='target', data=predicted_1, color='b',ax=ax);
9 ax.axvline(threshold,c='k');
```



Getting Predictions from sklearn

In [38]:

```
1 yhat_cl = logr.predict(X_proline)
2
3 predicted_0 = df_wine_2class[yhat_cl==0]
4 predicted_1 = df_wine_2class[yhat_cl==1]
5
6 fig,ax = plt.subplots(1,1,figsize=(8,6))
7 sns.scatterplot(x='proline',y='target', data=predicted_0, color='r',ax=ax);
8 sns.scatterplot(x='proline',y='target', data=predicted_1, color='b',ax=ax);
9 ax.axvline(threshold,c='k');
```



Getting Probabilities from sklearn

Getting Probabilities from sklearn

- said we could use output of logistic as $P(y = 1|x)$

Getting Probabilities from sklearn

- said we could use output of logistic as $P(y = 1|x)$

```
In [39]: 1 p_y_cl = logr.predict_proba(x_proline)
2 p_y_cl[:5].round(3) # p(y=0/x), p(y=1/x)
```

```
Out[39]: array([[0.982, 0.018],
 [0.977, 0.023],
 [0.997, 0.003],
 [1.    , 0.    ],
 [0.277, 0.723]])
```

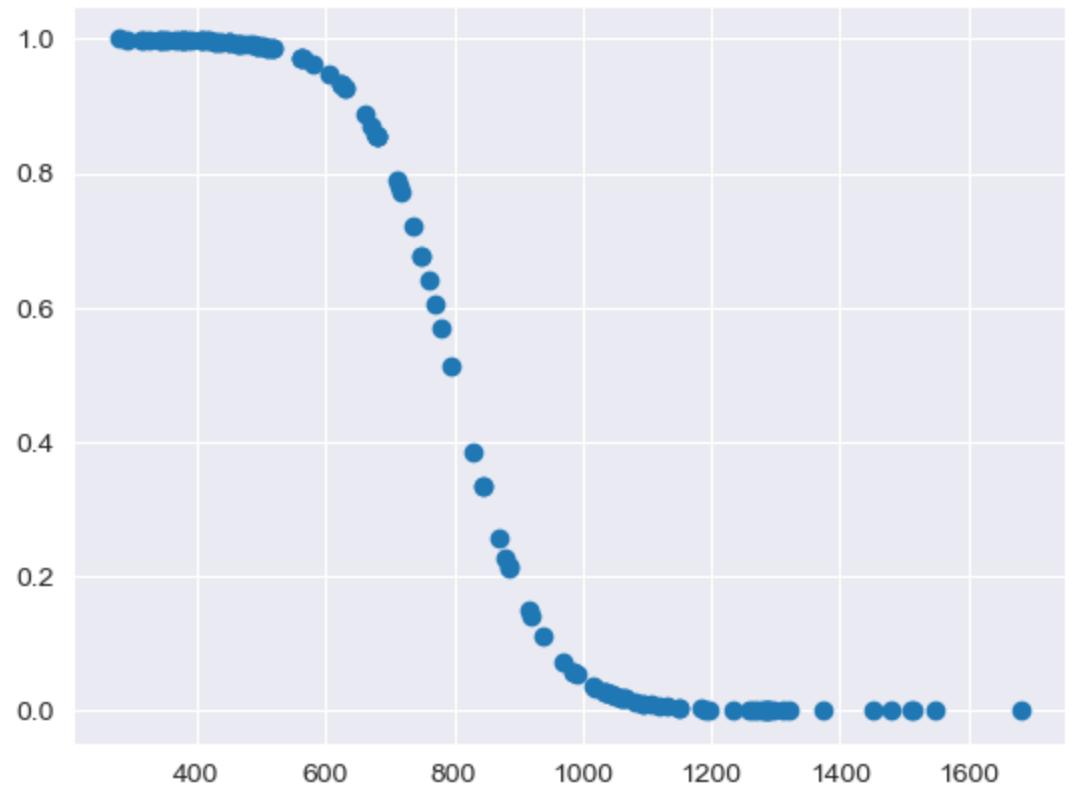
Getting Probabilities from sklearn

- said we could use output of logistic as $P(y = 1|x)$

```
In [39]: 1 p_y_cl = logr.predict_proba(X_proline)
2 p_y_cl[:5].round(3) # p(y=0/x), p(y=1/x)
```

```
Out[39]: array([[0.982, 0.018],
 [0.977, 0.023],
 [0.997, 0.003],
 [1.    , 0.    ],
 [0.277, 0.723]])
```

```
In [40]: 1 plt.scatter(X_proline,p_y_cl[:,1]);
```



Interpreting Logistic Regression Coefficients

- After some math

$$\log\left(\frac{y_i}{1-y_i}\right) = w_0 + w_1 x_{i1}$$

- this is the **log odds ratio** of $\frac{p(y=1)}{p(y=0)}$
- odds range from 0 to positive infinity
- odds(5) -> 5/1 -> 5 out of 6 times -> .83
- odds(.2) -> 1/5 -> 1 out of 6 times -> .16

See [here](#) for a good explanation

Logistic Regression with Multiple Features

Logistic Regression with Multiple Features

```
In [41]: 1 x_cl = df_wine_2class[['proline', 'hue']]  
2 x_cl_zscore = x_cl.apply(lambda x: (x-x.mean())/x.std())  
3  
4 logrm = LogisticRegression().fit(x_cl_zscore,y_cl)  
5  
6 for (name,coef) in zip(x_cl.columns,logrm.coef_[0]):  
7     print(f'{name:10s} : {coef: 0.3f}')
```

```
proline      : -3.464  
hue          :  0.488
```

Logistic Regression: plot decision regions

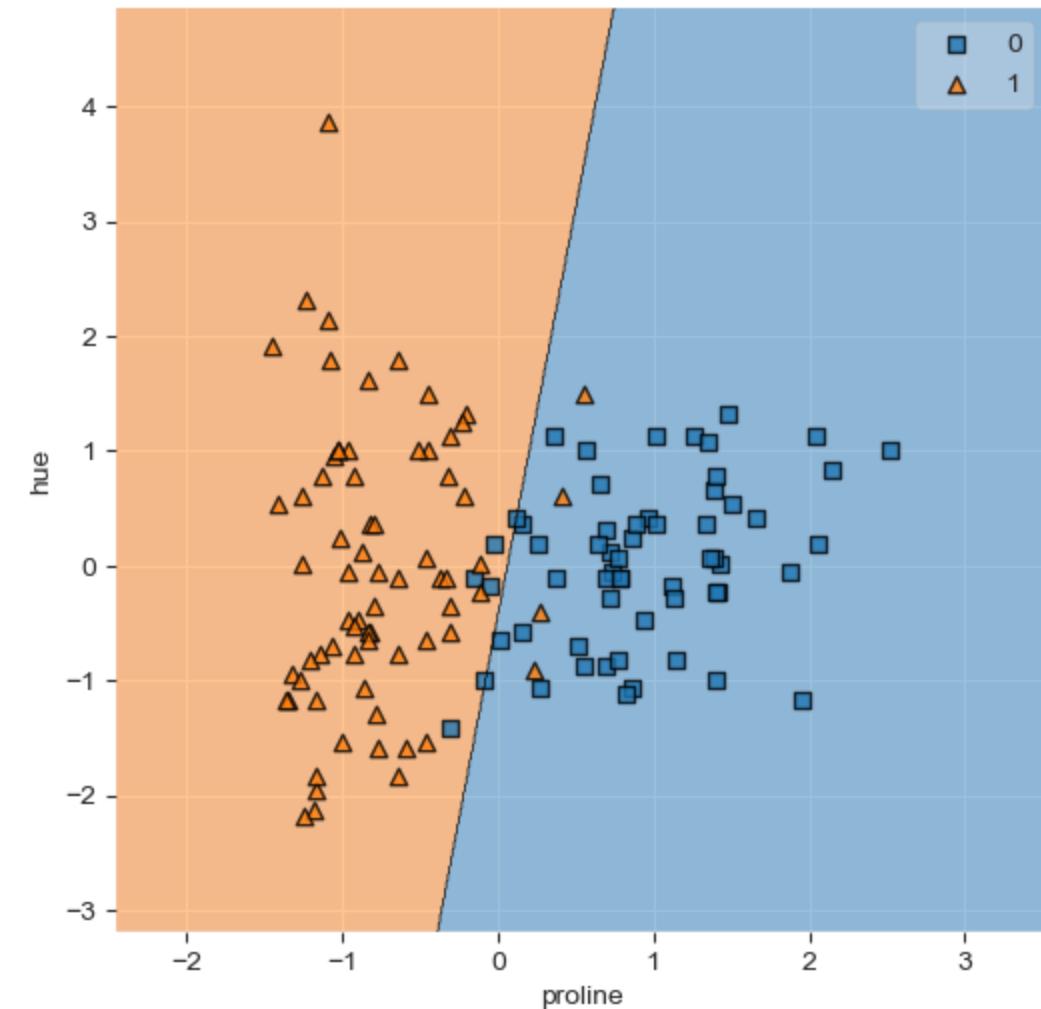
Logistic Regression: plot decision regions

```
In [42]: 1 print(f"X_cl.columns[0] : {logrm.coef_[0][0].round(3)} , X_cl.columns[1] : {logrm.coef_[0][1].round(3)}")  
proline : -3.464 , hue : 0.488
```

Logistic Regression: plot decision regions

```
In [42]: 1 print(f"X_cl.columns[0] : {logrm.coef_[0][0].round(3)}, X_cl.columns[1] : {logrm.coef_[0][1].round(3)}")  
  
proline : -3.464 , hue : 0.488
```

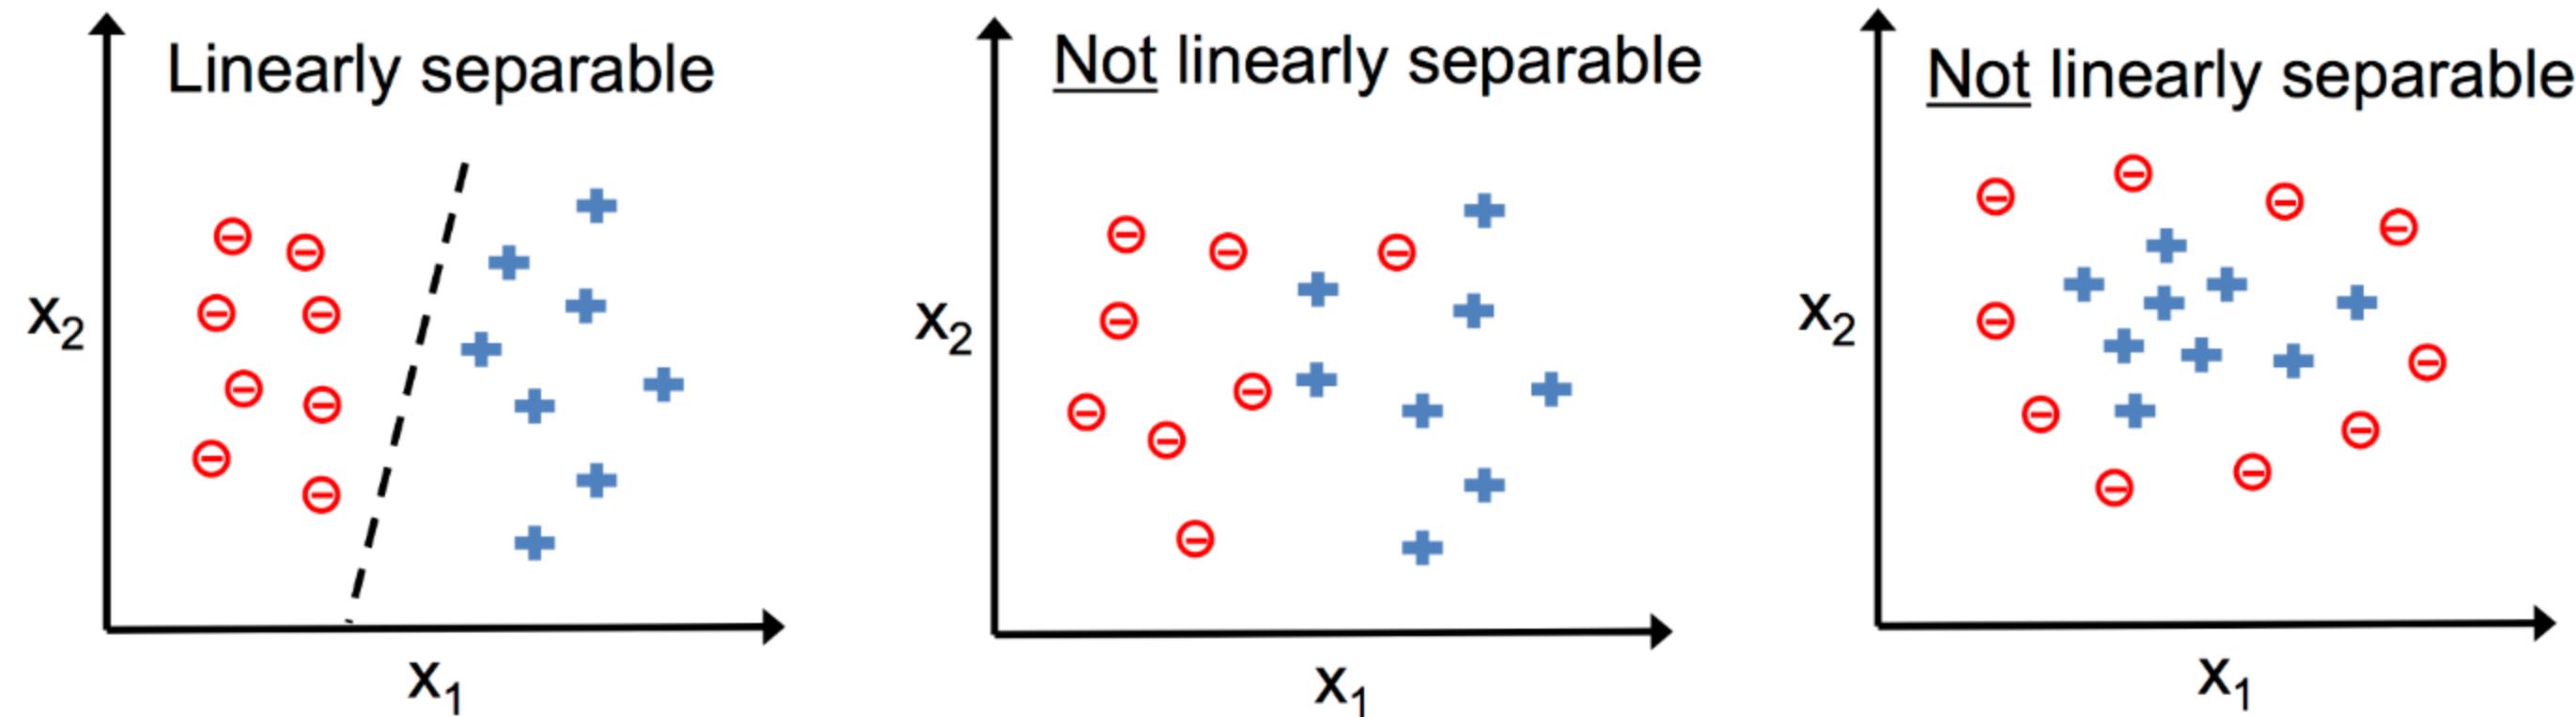
```
In [43]: 1 from mlxtend.plotting import plot_decision_regions # works with numpy arrays  
2  
3 logrm = LogisticRegression().fit(X_cl_zscore.values,y_cl.values)  
4 fig,ax = plt.subplots(1,1,figsize=(6,6))  
5 plot_decision_regions(X_cl_zscore.values, y_cl.values, clf=logrm, ax=ax);  
6 ax.set_xlabel(X_cl.columns[0]); ax.set_ylabel(X_cl.columns[1]);
```



Linearly Separable Data

Linearly Seperable Data

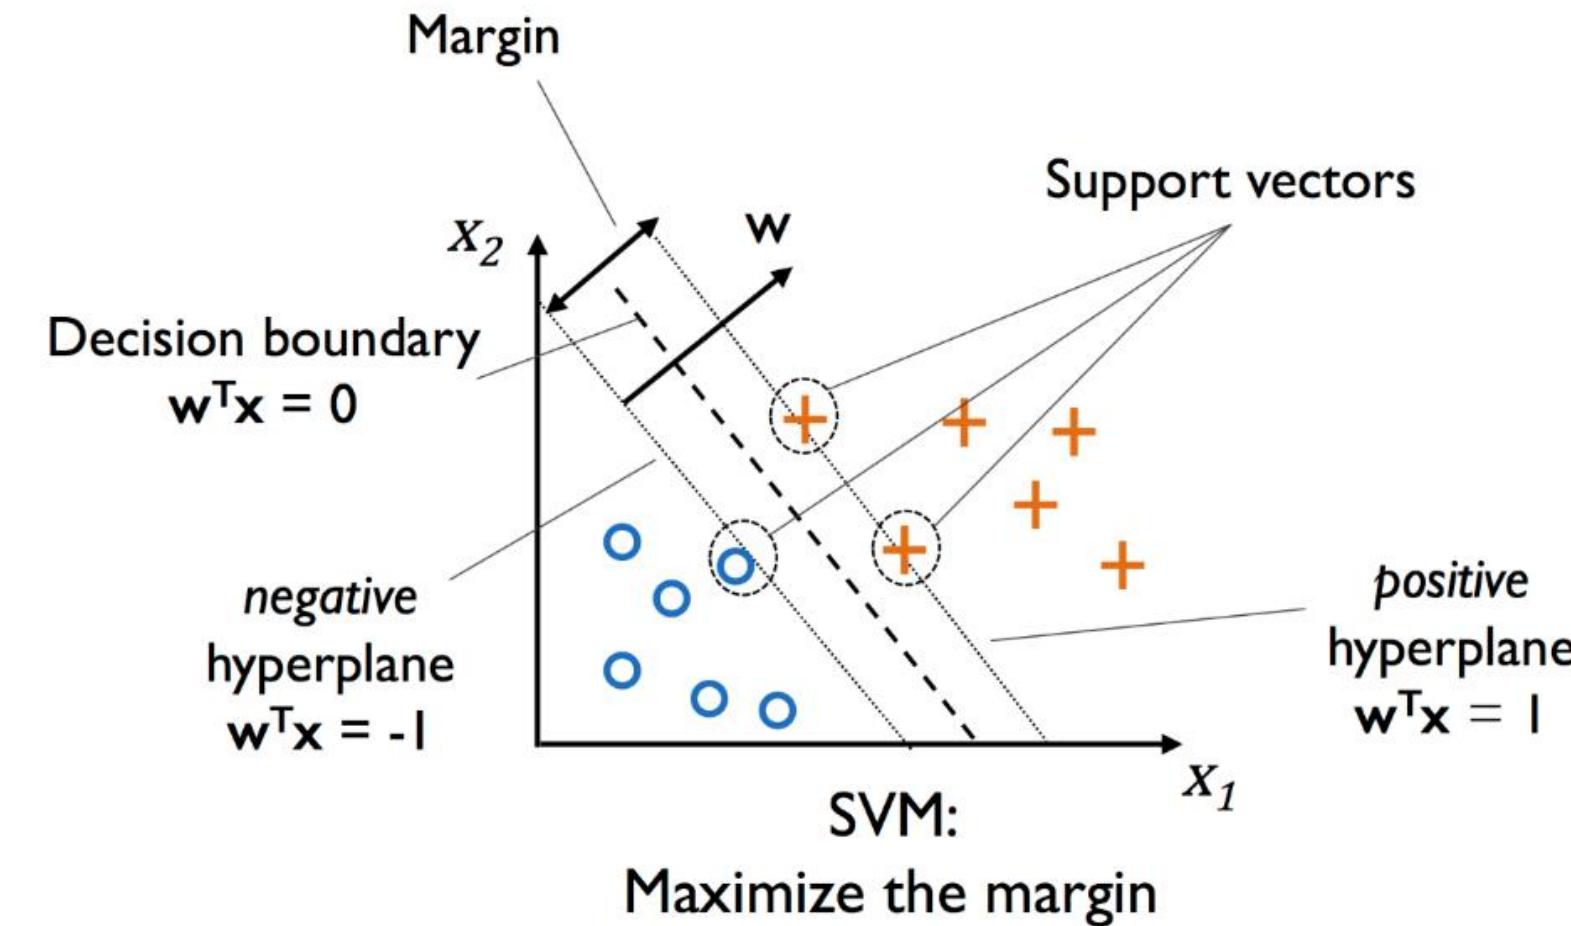
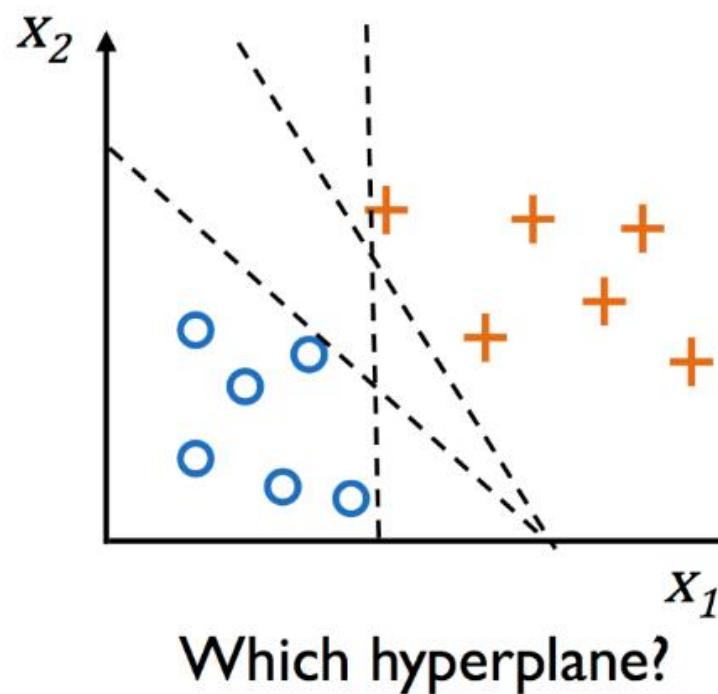
- Logistic Regression depends on data being linearly separable



From PML

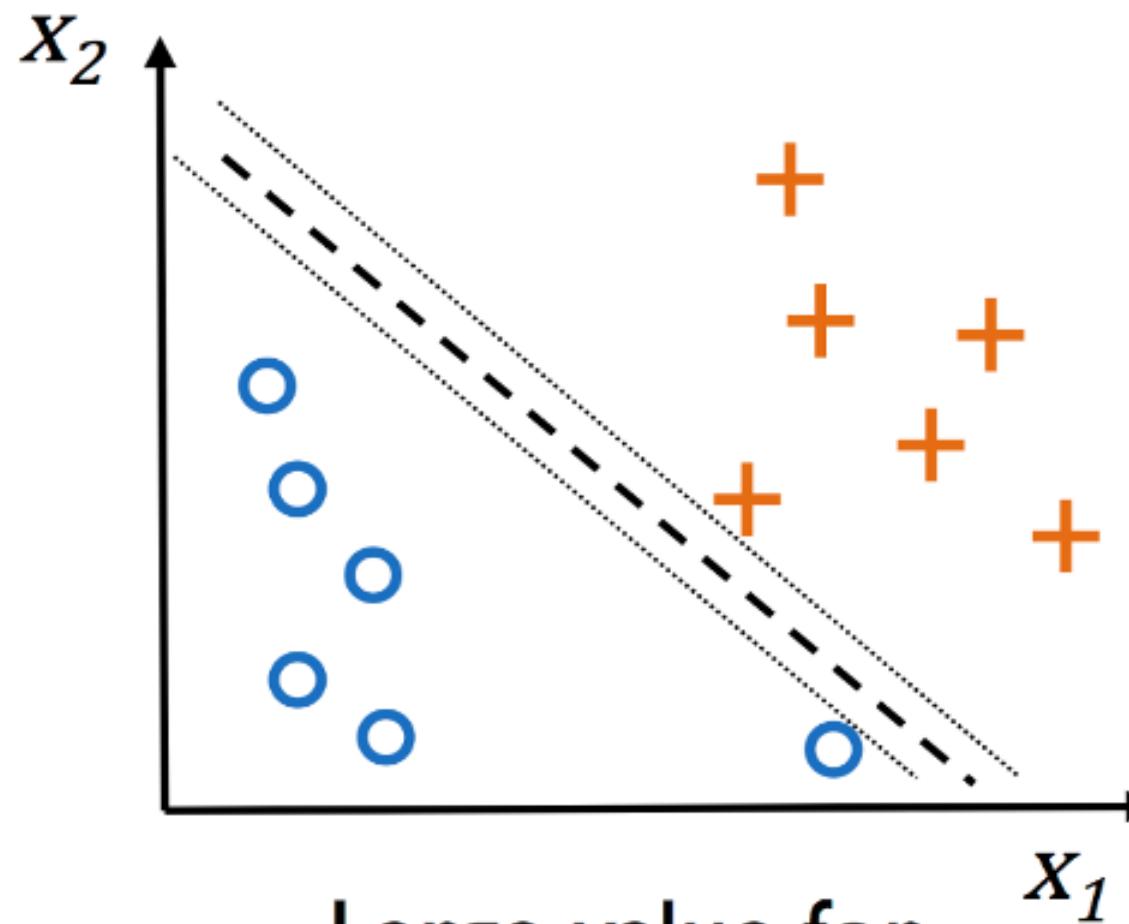
Which boundary to choose? Support Vector Machines (SVMs)

- For a linearly separable dataset, where should we place the decision boundary?
- Support Vector Machine (SVM) tries to "maximize the margin" between classes

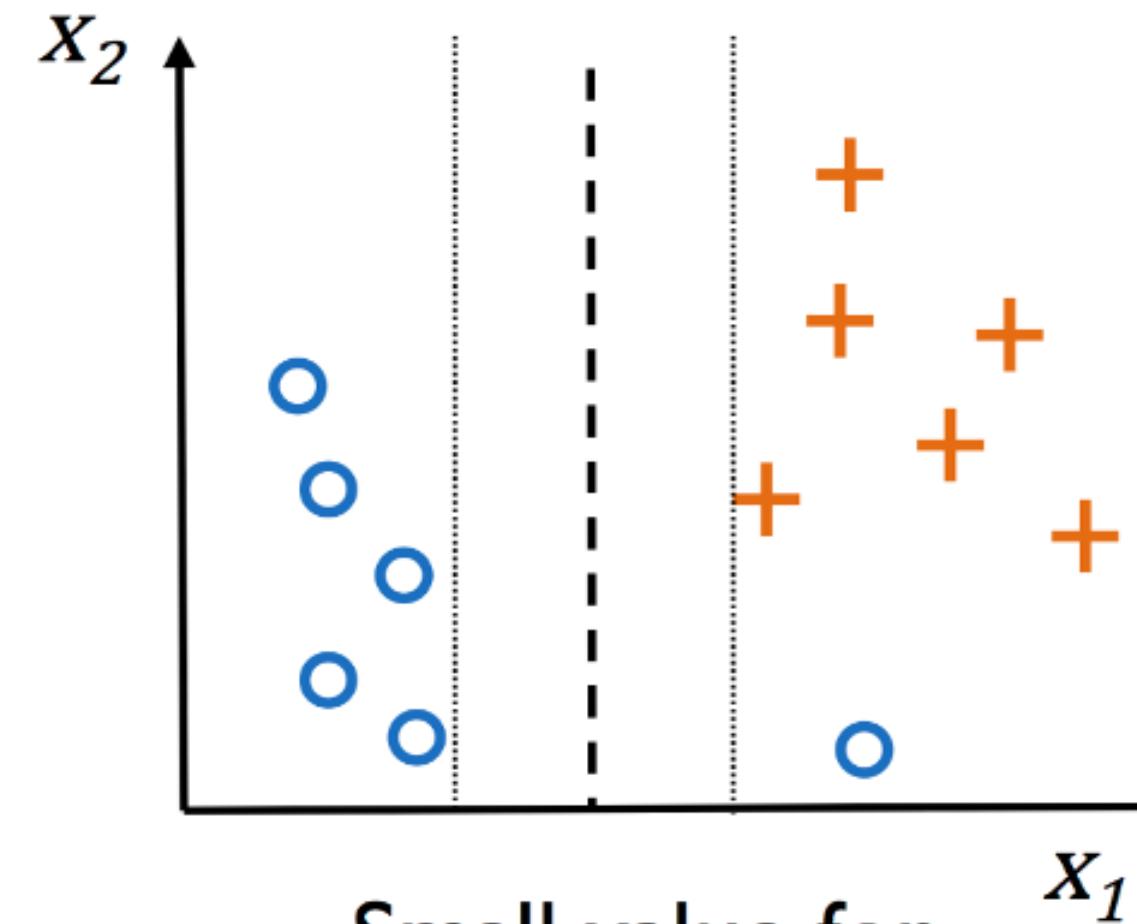


SVM Hyperparameter C

- Hyperparameter: Something we set



Large value for
parameter C



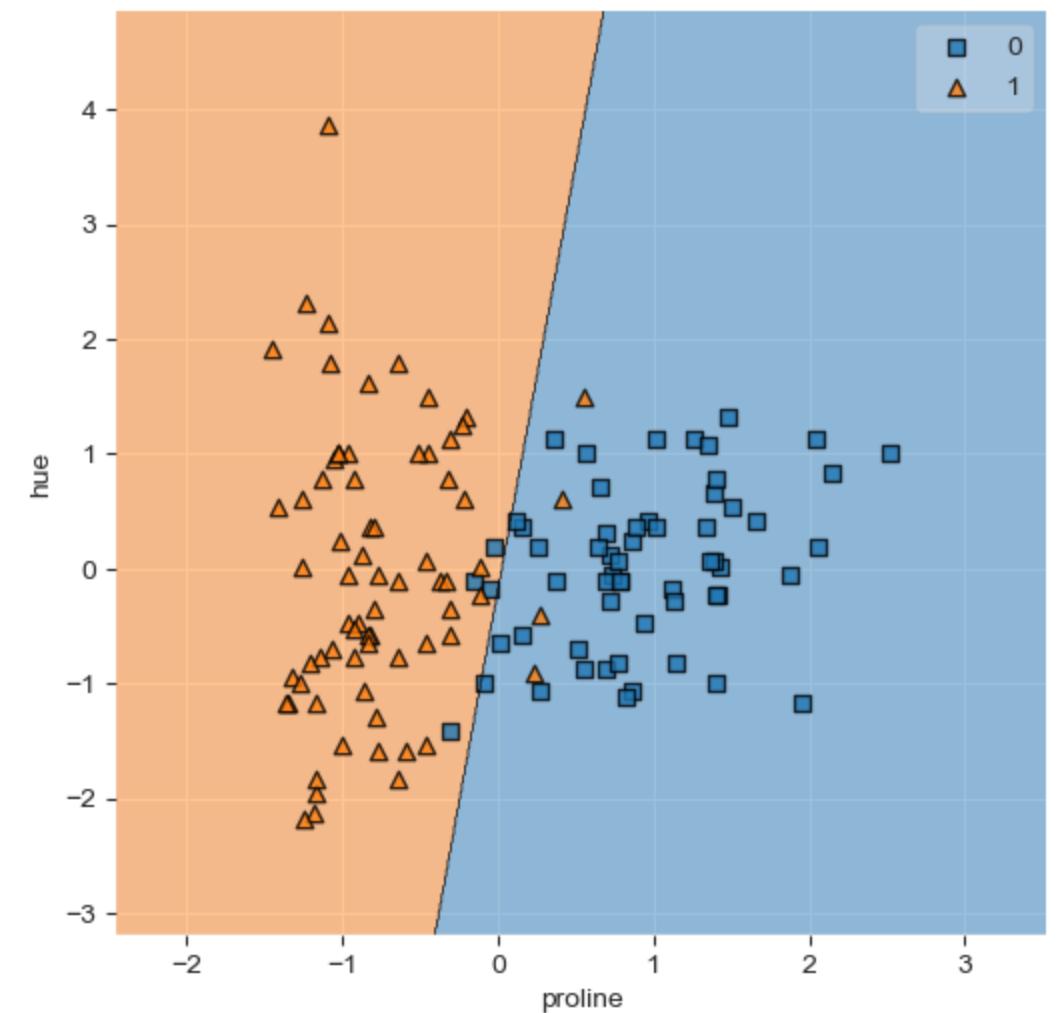
Small value for
parameter C

SVM with sklearn

SVM with sklearn

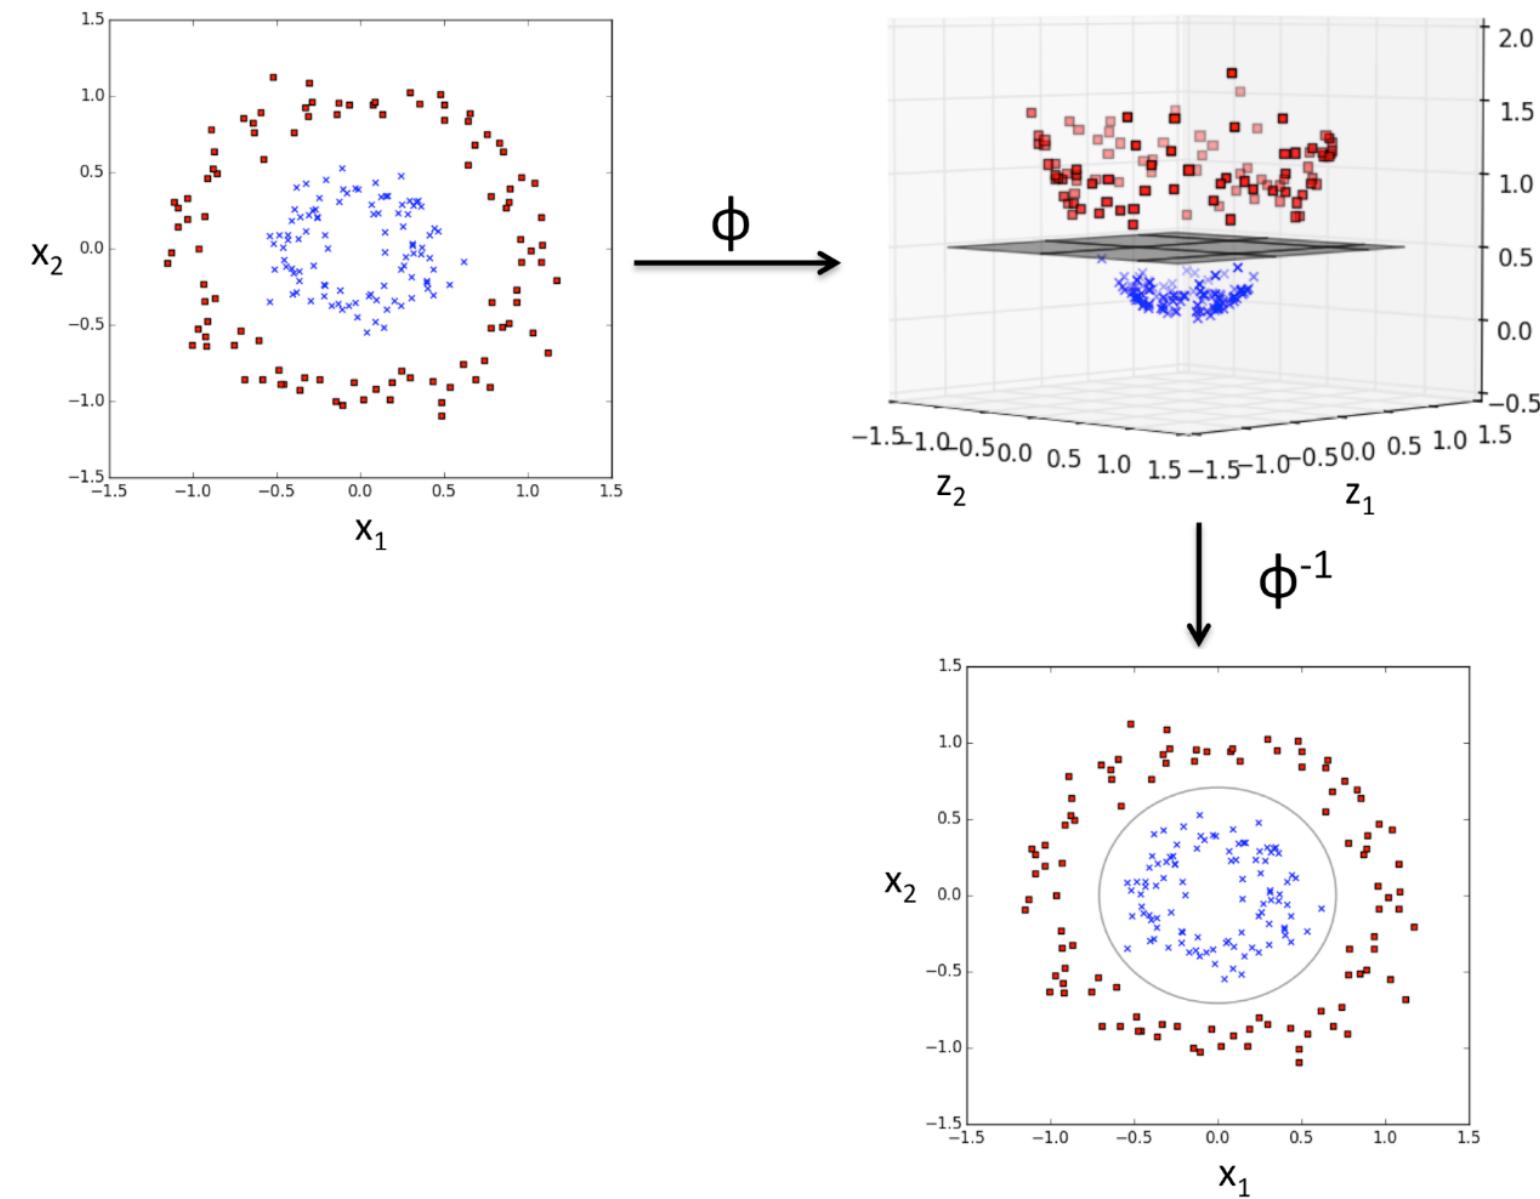
In [44]:

```
1 from sklearn.svm import SVC
2 svm_linear = SVC(kernel='linear')
3 svm_linear.fit(X_cl_zscore.values,y_cl.values);
4
5 fig,ax = plt.subplots(1,1,figsize=(6,6))
6 plot_decision_regions(X_cl_zscore.values, y_cl.values, clf=svm_linear);
7 plt.xlabel(X_cl.columns[0]); plt.ylabel(X_cl.columns[1]);
```



Non-Linear Boundaries with SVMs Kernel Trick

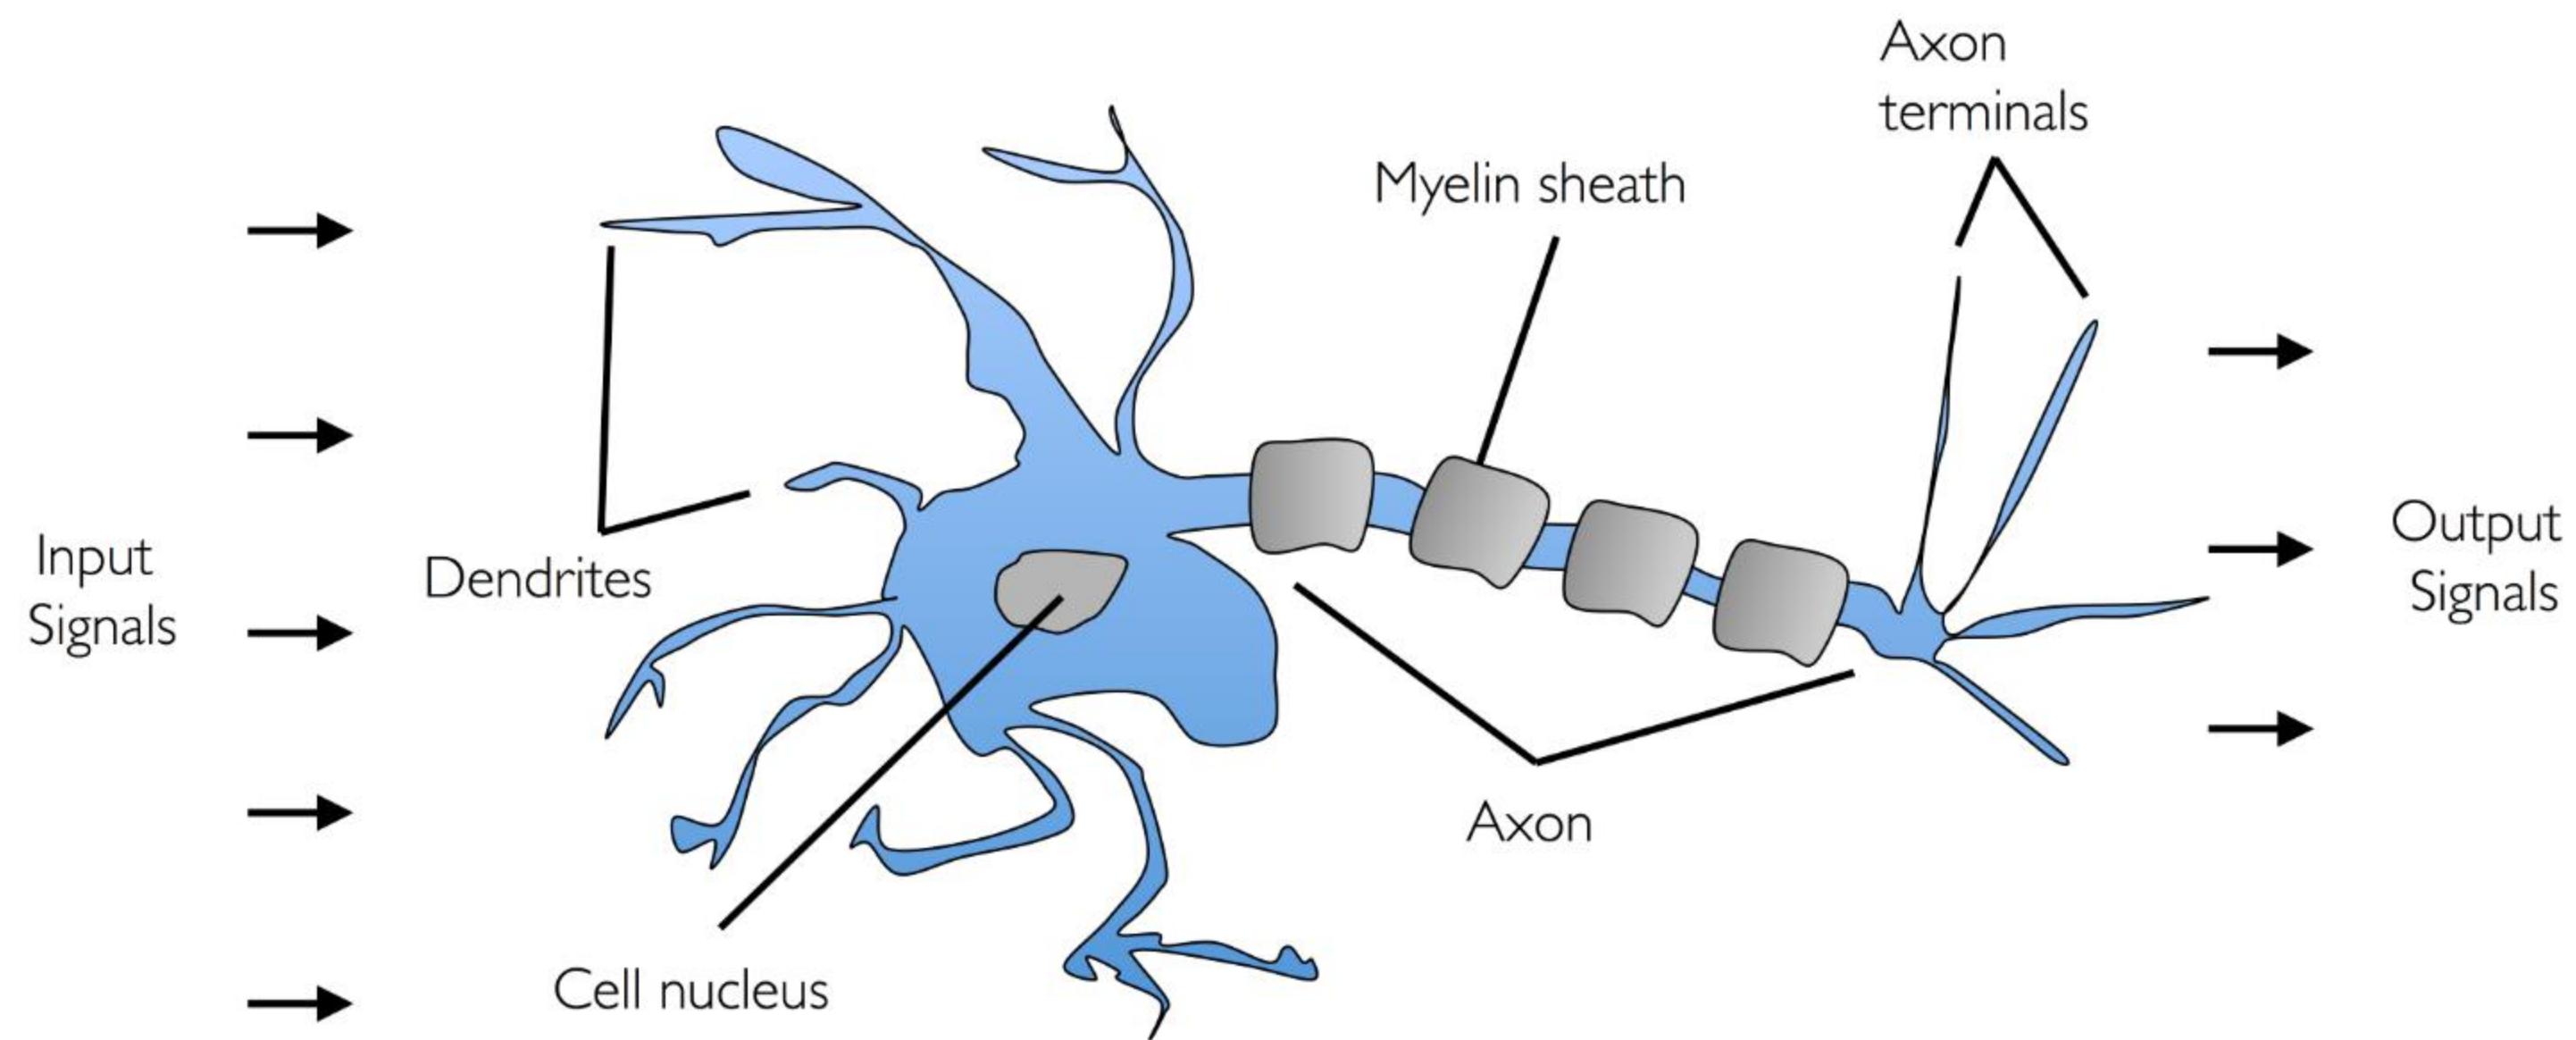
- Kernel Trick: Map data to a higher dimensional space and find linear boundary there



From PML

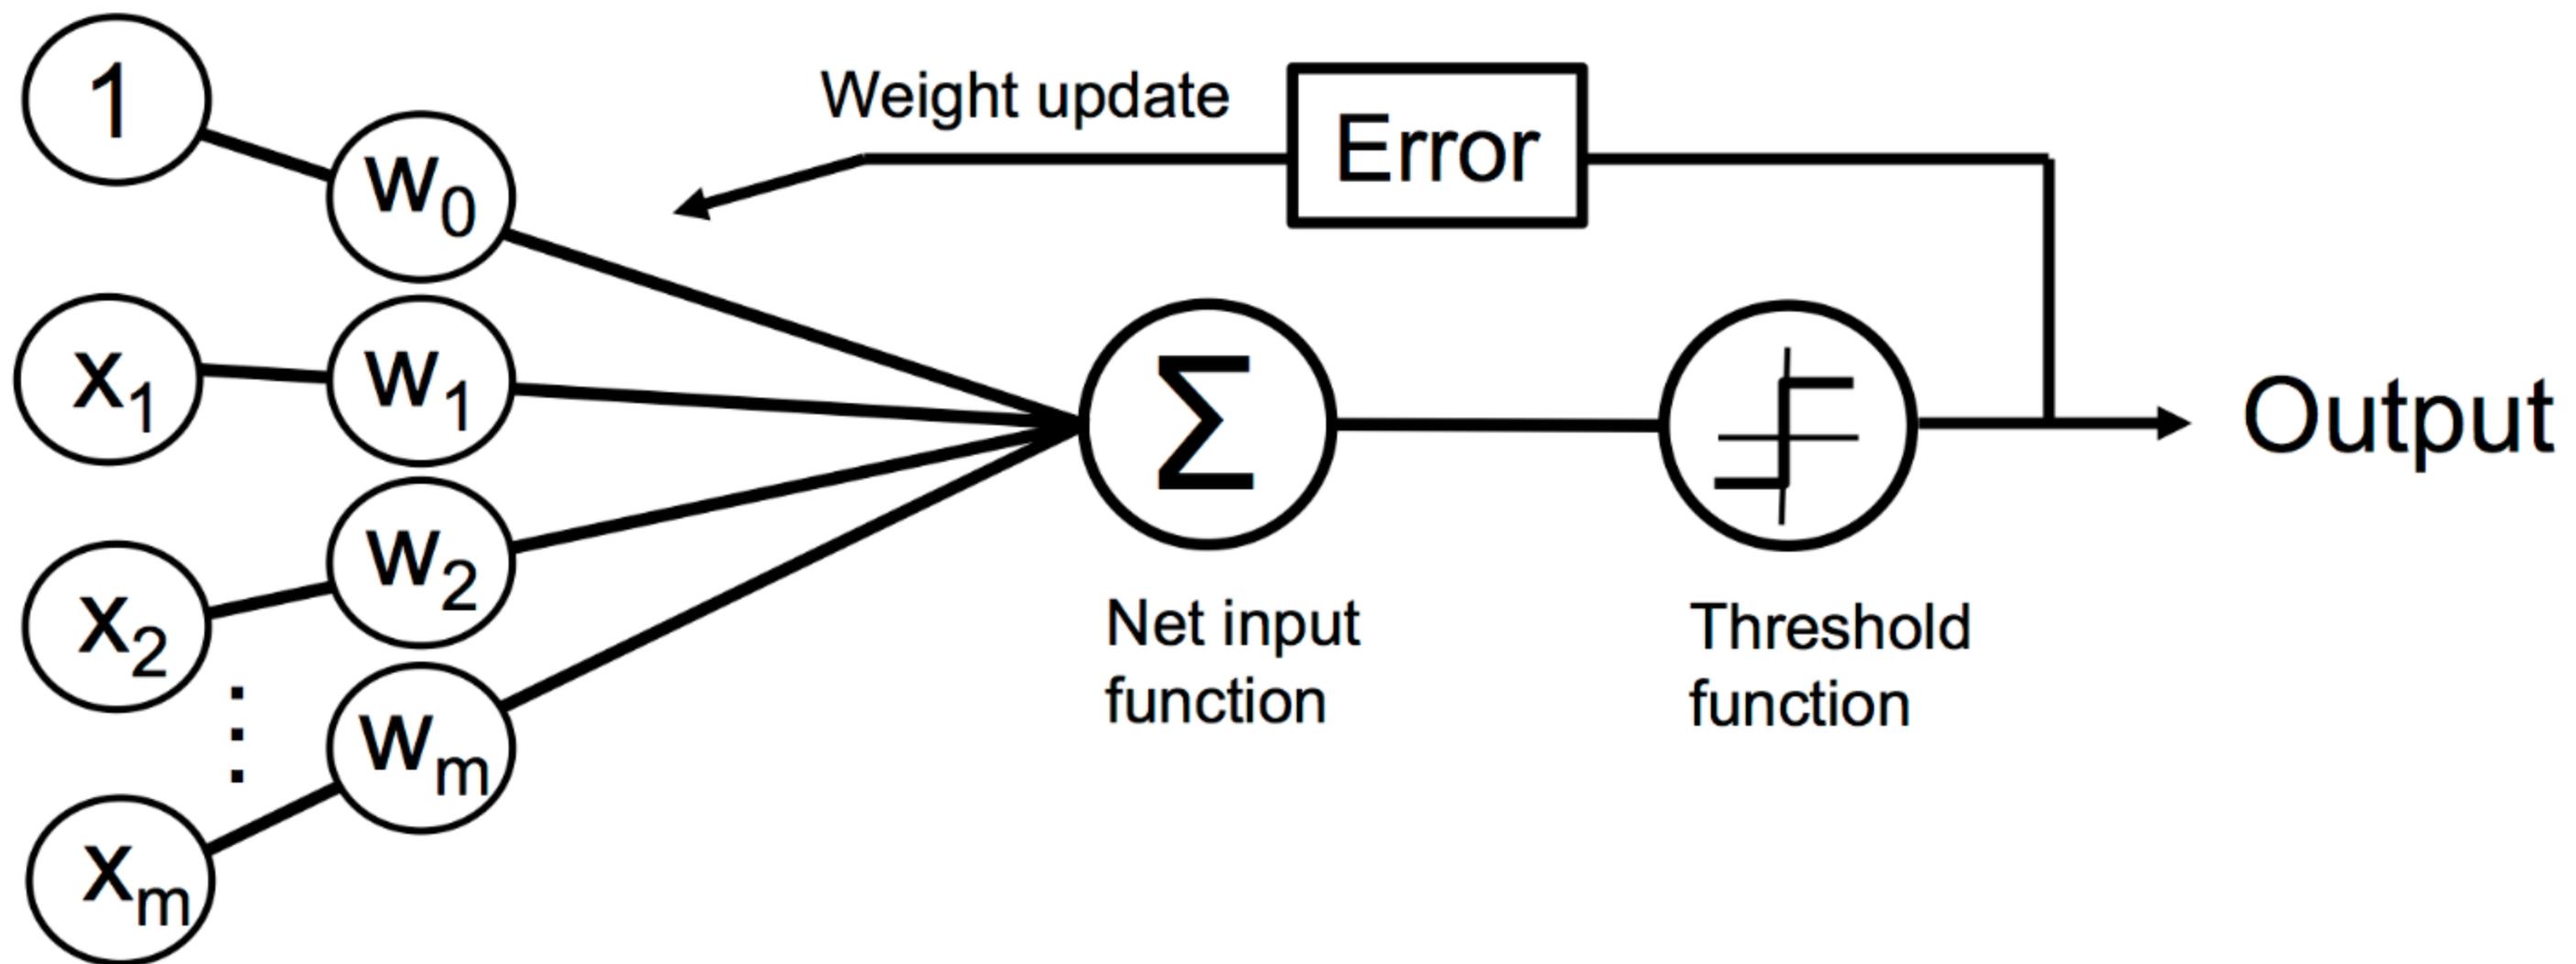
If we have time...

From Perceptron to Artificial Neural Network



From PML

Perceptron: Early Neuron Model



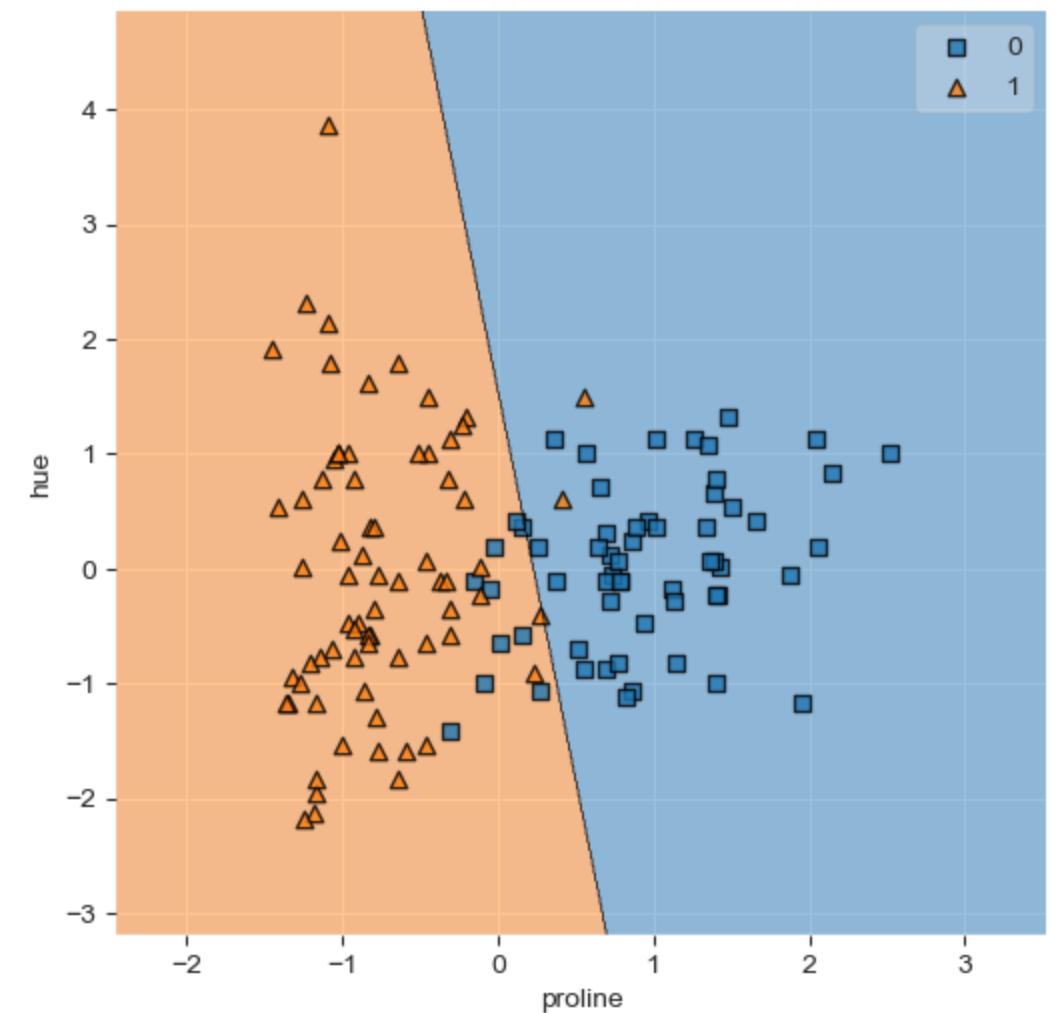
From PML

Perceptron in sklearn

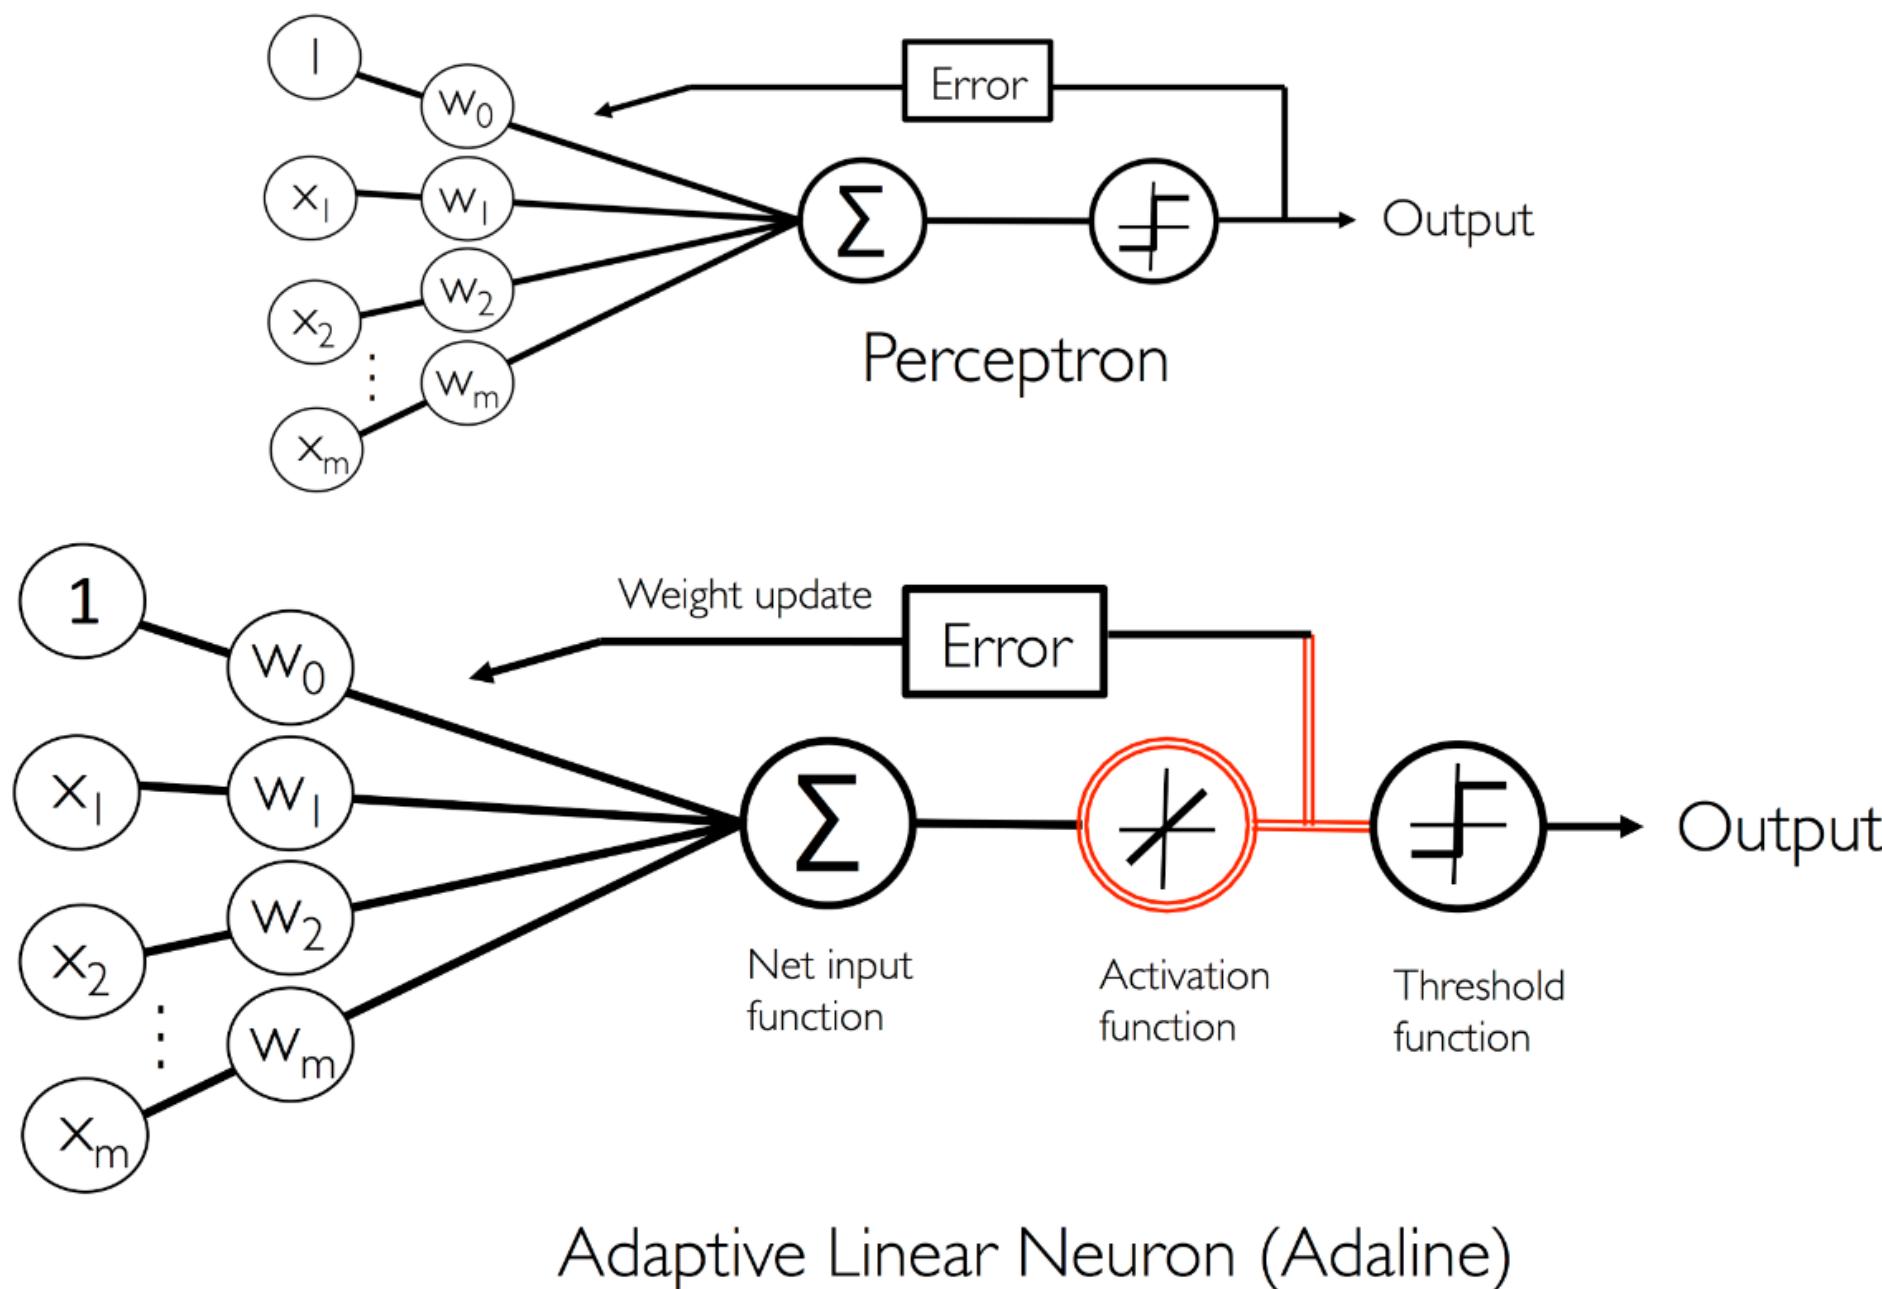
Perceptron in sklearn

In [46]:

```
1 from sklearn.linear_model import Perceptron  
2 perceptron = Perceptron()  
3 perceptron.fit(X_cl_zscore.values,y_cl.values);  
4  
5 fig,ax = plt.subplots(1,1,figsize=(6,6))  
6 plot_decision_regions(X_cl_zscore.values, y_cl.values, clf=perceptron);  
7 plt.xlabel(X_cl.columns[0]); plt.ylabel(X_cl.columns[1]);
```

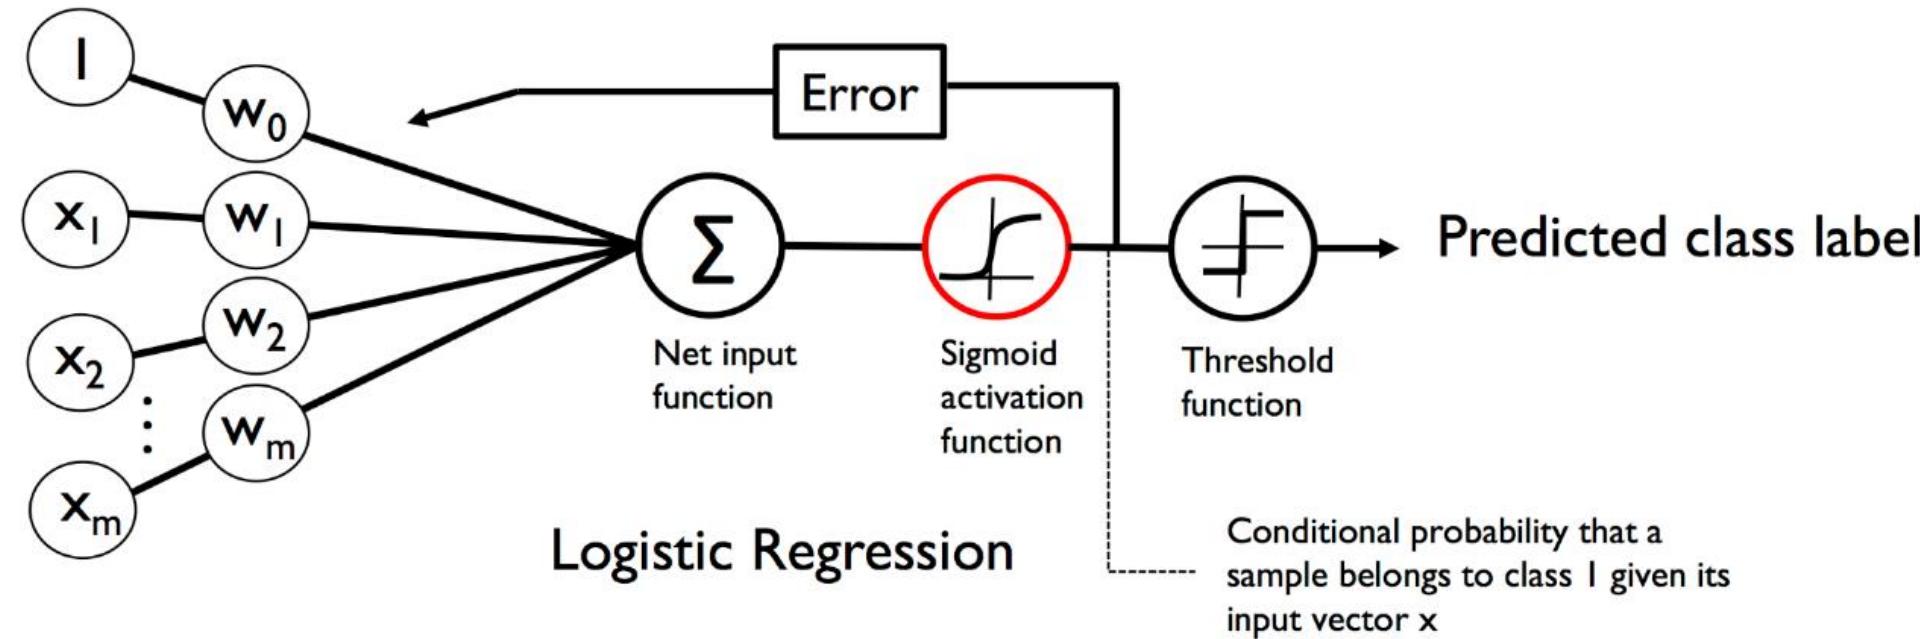
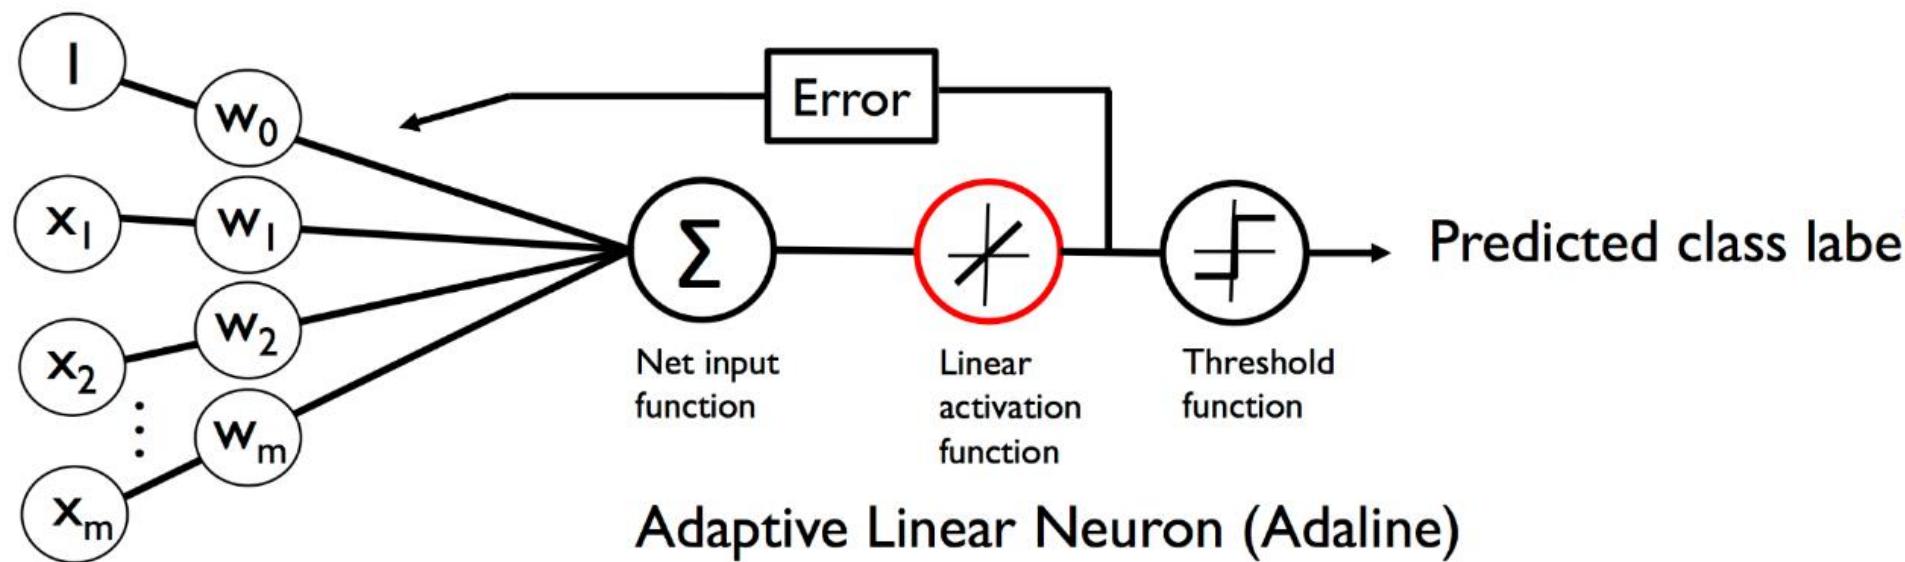


Perceptron to Adaline



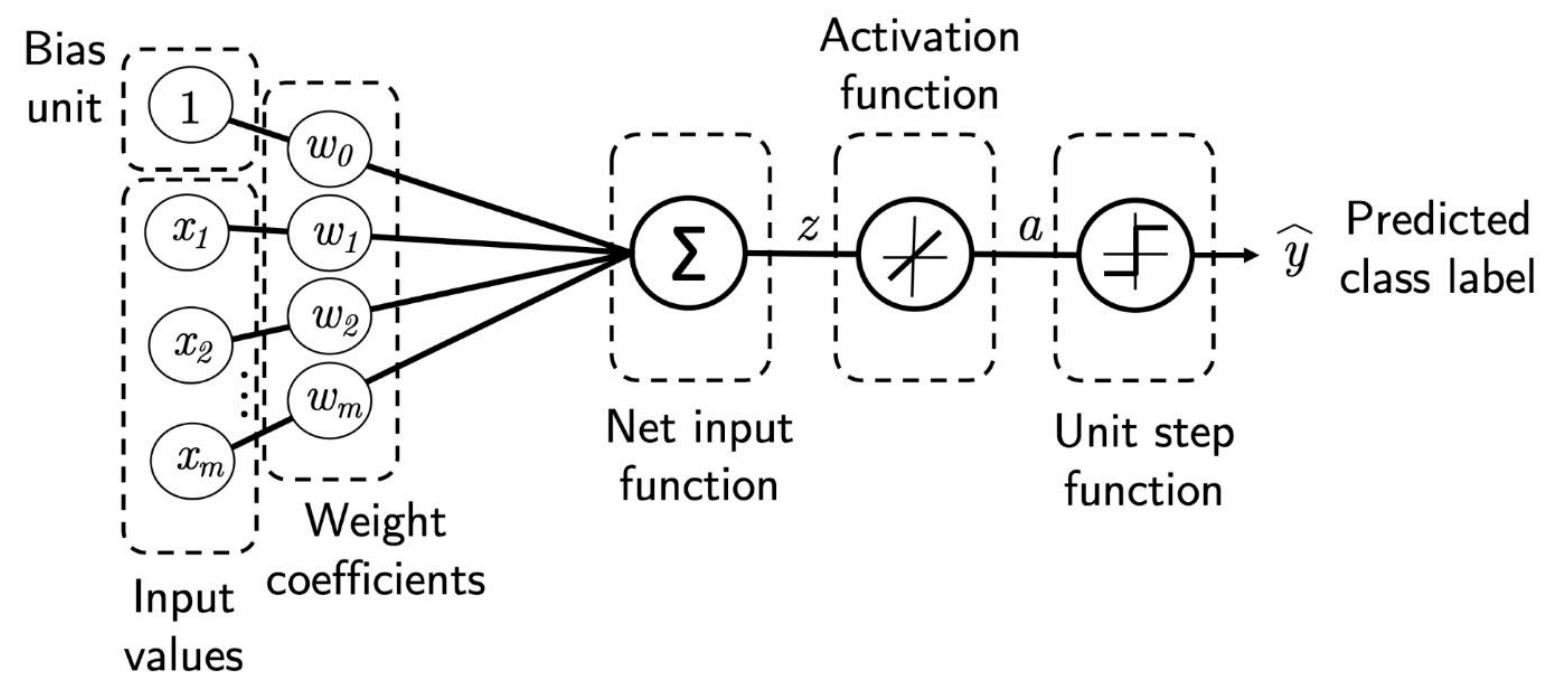
From PML

Adaline to Logistic Regression



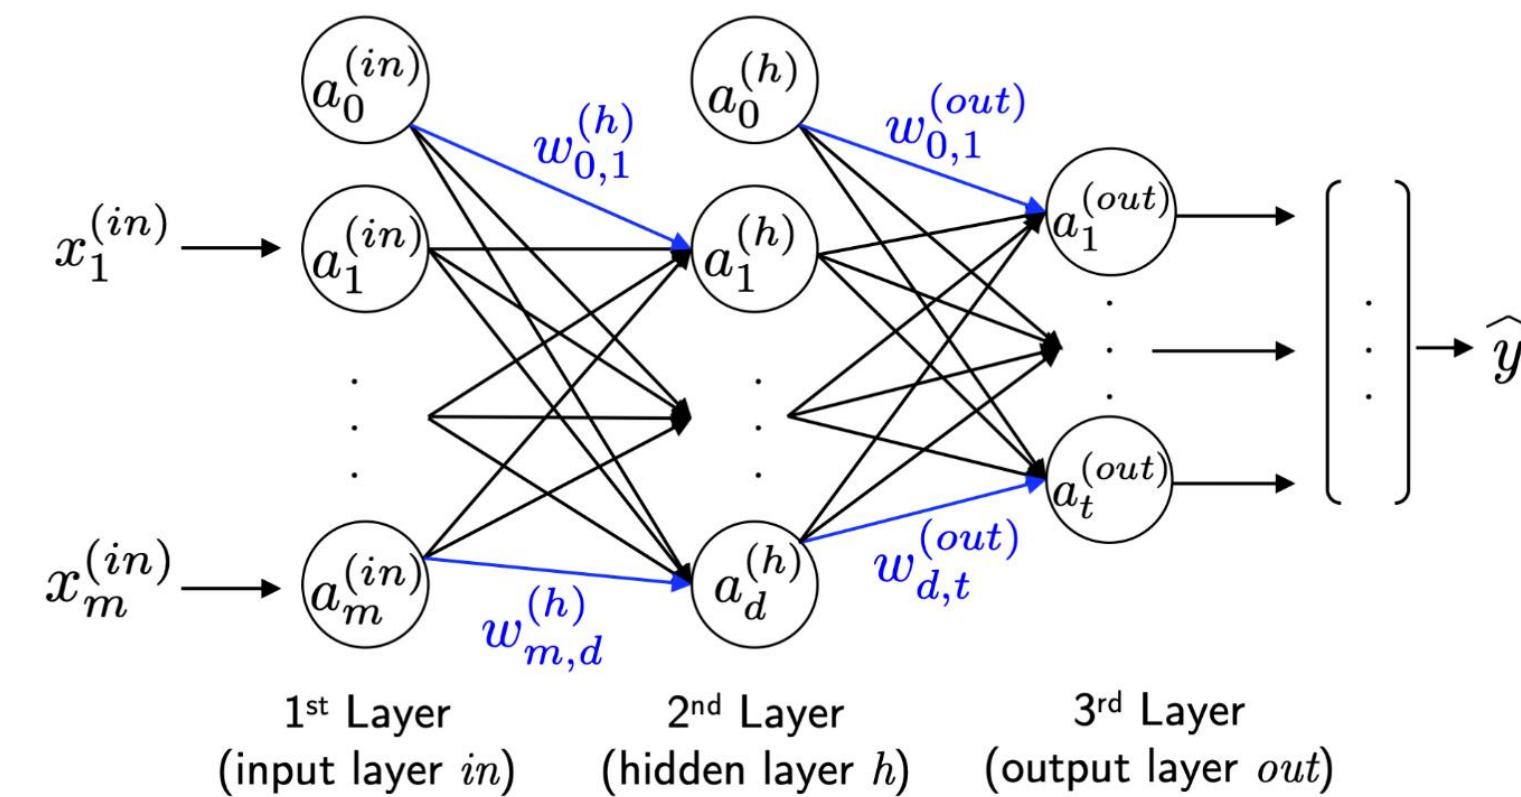
From PML

Components of Single Layer Neural Net



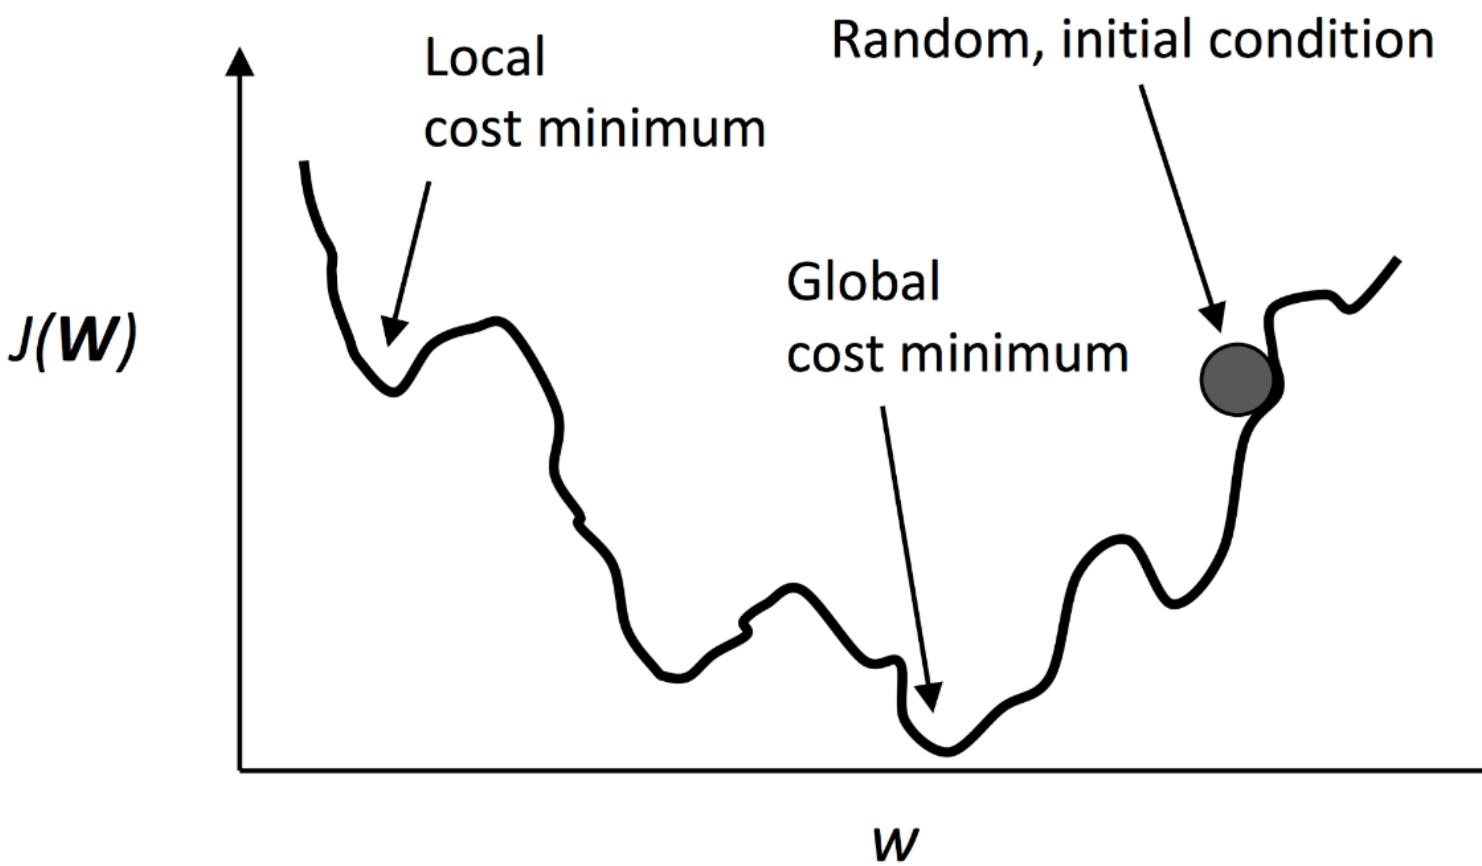
From PML

Multi-Layer Neural Network



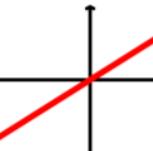
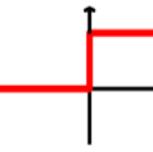
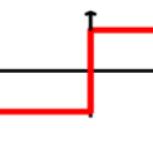
From PML

Complex Optimization Space



From PML

Activation Functions

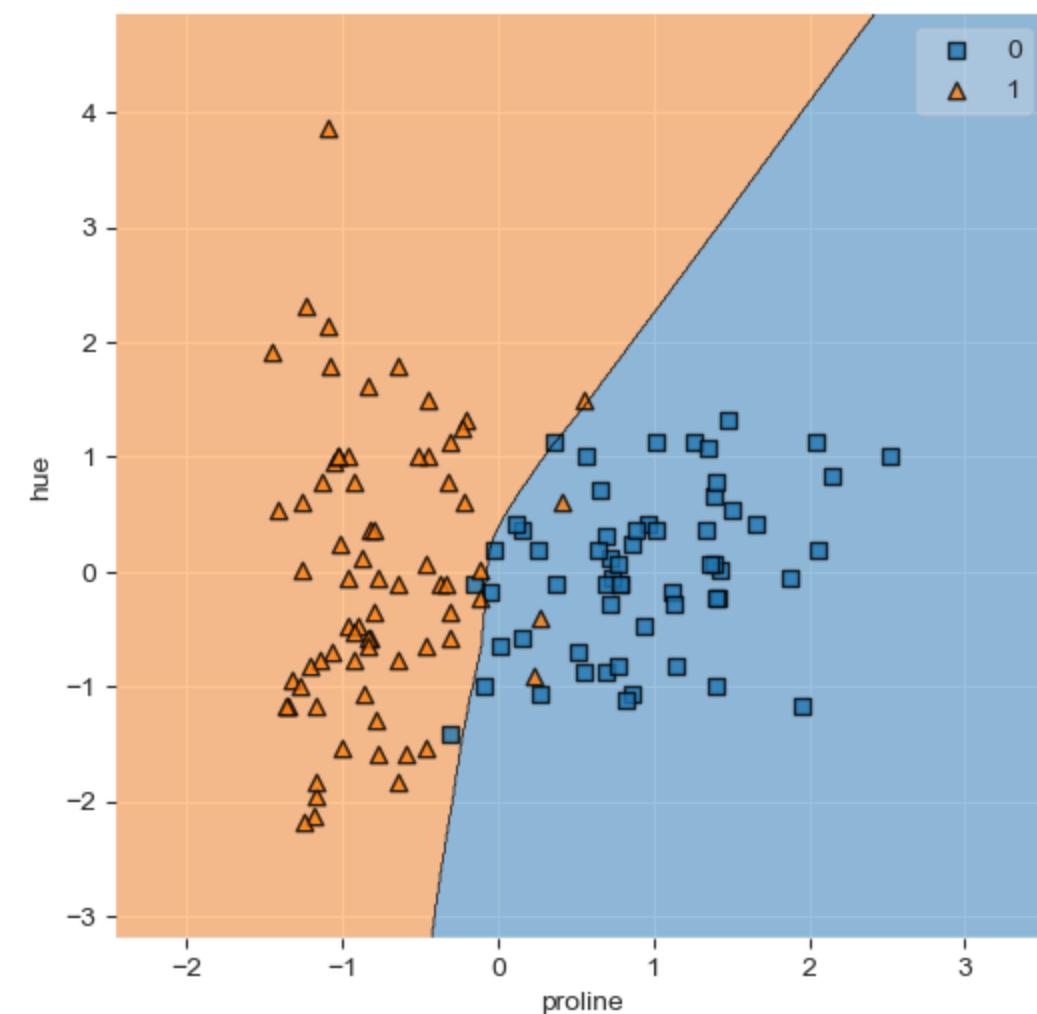
Activation function	Equation	Example	1D graph
Linear	$\phi(z) = z$	Adaline, linear regression	
Unit step (Heaviside function)	$\phi(z) = \begin{cases} 0 & z < 0 \\ 0.5 & z = 0 \\ 1 & z > 0 \end{cases}$	Perceptron variant	
Sign (signum)	$\phi(z) = \begin{cases} -1 & z < 0 \\ 0 & z = 0 \\ 1 & z > 0 \end{cases}$	Perceptron variant	
Piece-wise linear	$\phi(z) = \begin{cases} 0 & z \leq -\frac{1}{2} \\ z + \frac{1}{2} & -\frac{1}{2} \leq z \leq \frac{1}{2} \\ 1 & z \geq \frac{1}{2} \end{cases}$	Support vector machine	
Logistic (sigmoid)	$\phi(z) = \frac{1}{1 + e^{-z}}$	Logistic regression, multilayer NN	
Hyperbolic tangent (tanh)	$\phi(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$	Multilayer NN, RNNs	
ReLU	$\phi(z) = \begin{cases} 0 & z < 0 \\ z & z > 0 \end{cases}$	Multilayer NN, CNNs	

Multi-Layer Perceptron with sklearn

Multi-Layer Perceptron with sklearn

```
In [48]: 1 from sklearn.neural_network import MLPClassifier, MLPRegressor  
2 mlp = MLPClassifier(hidden_layer_sizes=(100,50,50,),  
3                      activation='relu',           # default  
4                      max_iter=100)  
5 mlp.fit(X_cl_zscore.values,y_cl.values);  
6  
7 fig,ax = plt.subplots(1,1,figsize=(6,6))  
8 plot_decision_regions(X_cl_zscore.values, y_cl.values, clf=mlp);  
9 plt.xlabel(X_cl.columns[0]); plt.ylabel(X_cl.columns[1]);
```

```
/Users/andi/miniconda3/envs/sigma2/lib/python3.8/site-packages/sklearn/neural_network/_multilayer_perceptron.py:702: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (100) reached and the optimization hasn't converged yet.  
warnings.warn(
```



Interpreting Multi-Layer Perceptron?

Interpreting Multi-Layer Perceptron?

```
In [49]: 1 len(mlp.coefs_)
```

```
Out[49]: 4
```

Interpreting Multi-Layer Perceptron?

```
In [49]: 1 len(mlp.coefs_)
```

```
Out[49]: 4
```

```
In [50]: 1 for w in range(len(mlp.coefs_)):
2     print(mlp.coefs_[w].shape)
```

```
(2, 100)
(100, 50)
(50, 50)
(50, 1)
```

Interpreting Multi-Layer Perceptron?

```
In [49]: 1 len(mlp.coefs_)
```

```
Out[49]: 4
```

```
In [50]: 1 for w in range(len(mlp.coefs_)):
2     print(mlp.coefs_[w].shape)
```

```
(2, 100)
(100, 50)
(50, 50)
(50, 1)
```

```
In [51]: 1 mlp.coefs_[0][0,:20].round(3)
```

```
Out[51]: array([ 0.178, -0.208, -0.231, -0.124, -0.237, -0.223, -0.193, -0.053,
 0.25 ,  0.257, -0.006, -0.273, -0.091, -0.15 , -0.066,  0.193,
-0.298,  0.248,  0.167, -0.176])
```

Interpreting Multi-Layer Perceptron?

```
In [49]: 1 len(mlp.coefs_)
```

```
Out[49]: 4
```

```
In [50]: 1 for w in range(len(mlp.coefs_)):
2     print(mlp.coefs_[w].shape)
```

```
(2, 100)
(100, 50)
(50, 50)
(50, 1)
```

```
In [51]: 1 mlp.coefs_[0][0,:20].round(3)
```

```
Out[51]: array([ 0.178, -0.208, -0.231, -0.124, -0.237, -0.223, -0.193, -0.053,
 0.25 ,  0.257, -0.006, -0.273, -0.091, -0.15 , -0.066,  0.193,
-0.298,  0.248,  0.167, -0.176])
```

- Further reading: [Interpretable Machine Learning: Model Agnostic Methods](#)
- Many packages for explainable/interpretable ML:
 - [InterpretML](#)
 - [SHAP](#)
 - [Shapash](#)
 - ...

Questions re Classification with Linear Models?